

Near-Total Weight-Space Erasure Without Behavioural Collapse

What Survives When a Merged LoRA Is Quantized

Technical report · Phase 0 and Phase 1

All numbers reproducible from the accompanying repository

`paper/appendix-D-reproduction.md`

Abstract

Merging a LoRA adapter and quantizing to INT4 at group size 128 delivers a weight update with **cosine 0.14** to the one the adapter trained, at **7.5 times its magnitude**. We lead with cosine because it does not depend on which tensor sets the quantization grid, where the fraction of weights that looks changed runs from 1.1% to 85.5% (§3.3).

Across **six rank-32 adapters sharing one base model and one training recipe**, quantized on the deployment path, the two sides of behaviour move differently. Elicitation *capability* shows **no detectable loss**: retention 99.2%, enumerated 95% interval [**90.7%, 107.6%**], spanning parity and excluding losses beyond about 9%. The trained *constraint* does move, and **tightens**: the adversarial leak rate falls from 16.7% to 8.3%, a paired difference of **+8.3 points** with an enumerated 95% interval of [**+4.2, +12.5**] — four of the six adapters leak one fewer prompt and none leaks more, on cells that are counts out of 8. An exact sign-flip permutation test gives **p = 0.125**: the direction is consistent and n=6 cannot resolve it (§5.1). §3.7 defines behaviour as two-sided, so a single "unchanged" would be wrong on both sides.

Both halves follow from one ratio, $|\Delta|/s$ — the adapter's per-weight magnitude against the quantization step. Read per weight, a channel model with **no fitted parameters** predicts each adapter's code-flip rate to within **2.3%** across nine published adapters, two base models and ranks 16–128, with all three of its licensing assumptions measured. Read on the adapter's r -dimensional active subspace, layer-output fidelity is *measured 6.2–16.5× higher* than weight-space fidelity — the span of B.13's per-adapter values, not a prediction from the concentration constant. Weight-space erasure and layer-output survival are one measurement read at two levels of the same map; behaviour is a third level and follows from neither (§5.3).

Capability degrades at coarser grids — 77.2% at INT4 per-channel, 57.8% at INT3 — and degrades **benignly**, the constraint holding throughout (§5.2). Weight-space measurement cannot say **which** adapter survives. Across six adapters matched on rank, scaling, base model, recipe and measured output SNR to within 3.3%, retention spans **28.7% to 86.4%** on the pre-registered instrument (28.4–84.4% floor-corrected), and six of the seven pairs the data can resolve run opposite to the predictor.

Notes for revision

- **Length.** ~280 words. Priority if trimmed further: keep the capability/constraint split with both intervals, the channel model's 2.3%, and the predictive gap.
- The opening contrast mirrors **Figure 1** (twin panel) by design.
- "99.2% retention" is elicitation score relative to the same adapter's own BF16 score — stated precisely in §5.1. The abstract must not drift into implying a broader behavioural claim than §8.1 licenses, and must not restate 99.2% without the interval: the interval is what makes it a bounded non-detection rather than a measured equality.
- **The capability claim may not be stated as a claim about behaviour.** An earlier version read "no detectable change in their trained behaviour", which §3.7's own two-sided definition

falsifies: the constraint side is measured and it moves. The non-detection is on capability alone, and the constraint result goes beside it.

1. Introduction

The contrast

Take a published LoRA fine-tune, merge it into its base model, and quantize the result to 4 bits with group size 128 — the ordinary deployment path. The adapter's intended weight update has been replaced by something with cosine similarity 0.14 to it, and 7.5 times its magnitude, pointing in a direction the adapter never requested. Not one of that adapter's weight deltas reaches even a quarter of a quantization half-step; the median falls below the step size by a factor of 128.

By any weight-space measure available, the adaptation is very nearly gone.

How much of the model looks changed depends on which tensor sets the quantization grid and on which count you take (§3.3), and the readings run from 1.1% to 85.5%. On the deployment path a single adapter shows **85.5% of stored values** changed against **2.1% of the integer codes**, a factor of 41 for that adapter and 15.0× pooled over the nine; holding the grid fixed puts its value changes at 1.1%. Those are two different contrasts — one between regimes, one between metrics — and an earlier draft quoted one adapter's pair with the pooled ratio, which is how 41 came to be printed as 15.

This is why the headline of this paper is a cosine and not a count. The cosine between the intended and delivered update moves from 0.1390 to 0.1379 between the regimes, a 0.8% difference. The question the paper asks — did the trained update survive? — has a regime-independent answer; the question the counts answer — how much of this checkpoint differs from the base model? — does not.

On the same models, quantized the same way, elicitation capability is indistinguishable. Retention is 99.2%, and its enumerated 95% interval — [90.7%, 107.6%] — spans parity. The instrument cannot separate the quantized model from the unquantized one. What it does establish is a bound: **losses beyond about 9% are excluded.**

The other half of behaviour does move. §3.7 defines the battery as two-sided — capability and constraint — and the constraint **tightens**: the same six adapters leak the suppressed word on 16.7% of adversarial prompts at BF16 and 8.3% after quantization, a paired **+8.3 points** with an enumerated 95% interval of **[+4.2, +12.5]** that excludes zero (§5.1). So "no behavioural change" is not what was measured and is not what this paper claims: one side is unchanged and the other moves, in the direction of disclosing less.

Figure 1 puts the weight-space and capability panels side by side. This paper is about that contrast: how both statements can be true at once, what governs each of them, and what follows for anyone shipping a quantized fine-tune.

Why this matters, and why the obvious worry is the wrong one

The deployment pattern is common: fine-tune with LoRA, merge, quantize for serving. If a rank-16 adapter produces a weight delta smaller than the 4-bit step size, the merged model is numerically

indistinguishable from the base model — and an "aligned quantized model" could in principle be behaviourally the base model with the alignment silently absent. That concern is what motivated this work, and it is asserted in adjacent literature: work on quantization-robust unlearning states that ordinary fine-tuning updates are "too small to survive 4-bit quantization" and that quantization can cause models to revert to pre-unlearning behaviour (§2.4).

We find the weight-space half of that concern to be correct and the behavioural half not to follow. The stored weights really are almost unchanged. The behaviour is not. The reason is not that the weight measurement is wrong — we validate our quantizer bit-exact against `gptqmodel` on 36 of 36 configurations — but that weight-space fidelity is the wrong quantity for predicting behaviour, in a way that is derivable rather than accidental.

The governing quantity

A single ratio explains both halves: $|\Delta|/s$, the adapter's per-weight magnitude against the step size of its quantization group.

At the level of stored weights, $|\Delta|/s$ determines everything. Treating quantization as a stochastic rounding channel gives a code-flip rate of $\text{mean}(\min(|\Delta|/s, 1))$, and this closed form — with **no fitted parameters** — predicts the measured flip rate of every adapter we test to within **2.3%**, across two base models, ranks from 16 to 128, both scaling conventions, and four different training regimes. What licenses the form is a measured property rather than goodness of fit: trained deltas are statistically independent of where a weight sits within its quantization bin (correlation < 0.0011), because gradient descent optimises a loss with no knowledge of the deployment quantizer.

At the level of layer outputs, the same ratio behaves differently, because a rank- r adapter concentrates its effect on an r -dimensional subspace while quantization error spreads across all d_{in} input directions. On inputs inside that subspace, signal is amplified relative to noise by $\sqrt{d_{\text{in}}/r}$ — a factor of 6.2–16.5 across the nine adapters measured, at ranks 16 to 128. A layer whose weight-space cosine is 0.14 can carry an output signal-to-noise ratio near 1.6. **Weight-space erasure and layer-output survival are not in tension; they are one measurement read at two levels of the weight-to-output map.** Behaviour is a third level and follows from neither.

What we find

- 1. Erasure with survival.** At INT4 g128 on the deployment path, stored values change in 85.5% of positions and stored codes in 2.1%, with 98.9% of codes unchanged when the grid is held fixed, and **no loss of elicitation capability is detectable** — while the constraint side tightens by 8.3 points — **measured on the same six adapters under the same regime.** The mean degrades as the grid coarsens — 99.2% $\rightarrow$ 77.2% $\rightarrow$ 57.8% across INT4 g128, INT4 per-channel and INT3 g128, all three paired contrasts excluding zero (B.8) — but **individual adapters do not**: four of six fall monotonically across the three quantized grids and only `gold` falls at every step including BF16, and at INT3 the mean sits between two adapters below half and two above 80% on the pre-registered instrument (three and one floor-corrected, B.7).
- 2. A parameter-free channel model** predicting weight-space retention within 2.3% across every adapter tested, together with a derived subspace-amplification law accounting for why behaviour outlives weights.

3. **The dissociation is benign.** Where behaviour does degrade, the trained *capability* weakens while the trained *constraint* holds (suppression ratio 0.18–0.27, Cliff's *d* between –0.56 and –0.83 at every precision). This is the opposite of the alarming failure mode — retained knowledge with lost restraint — and it is the opposite of what we predicted before measuring it.
4. **The predictive gap.** Weight-space retention does not predict behavioural retention. Within six adapters matched on rank, scaling, base model, recipe *and* measured output SNR to 3.3%, behavioural retention spans 28.7%–86.4% (28.4–84.4% floor-corrected); six of the seven pairs the data can resolve invert the ordering; and the adapter with the largest weight-space footprint has no measurable target behaviour at all. This limits what our own released tool can claim, and we say so in the tool's output.

Contributions

- **A parameter-free model of adapter retention under merge-then-quantize**, validated within 2.3% on nine published adapters spanning 1.1 decades of magnitude, and on a synthetic sweep across three decades, with all three of its licensing assumptions measured rather than assumed (§4.1), the first of them conditionally. The third was created by an argument we added a revision round after the model, and counted only when a reader counted for us.
- **A derived subspace-amplification law**, $\sqrt{(d_{in}/r)/(1+c/r)}$, reconciling weight-space erasure with *layer-output* survival. Its single empirical input $c \approx 0.87$ is a correction term predicted by the channel model, not a fitted scale (§3.6, §4.4). It accounts for the **level** — why a weight-space SNR of 0.13 still leaves subspace signal above noise — and, by §5.4, not for which adapter beats which.
- **End-to-end behavioural measurement** on the same models as the weight measurement, establishing no detectable loss at INT4 g128, a dose-response monotone in the mean but not in every adapter, and the benign direction of the dissociation (§5.1–5.3).
- **The predictive gap** (§5.4): three demonstrations, differing in kind, that weight-space measurement does not predict behavioural outcomes.
- **A weight-space account of why two opposing results can both hold.** Prior work reports that compressing *delta weights* protects alignment; we report that merged adapters are weight-space-erased. We argue these are the same law evaluated at opposite ends of $|\Delta|/s$, distinguished by **which tensor sets the quantization scale** (§2.5), and we measure the unmerged case rather than arguing it (`METHODOLOGY.md`). Their claim is behavioural and ours is not, so this is an account of the mechanism, not a reconciliation of findings.
- **ar.predict**, a tool computing effective adapter magnitude from published checkpoints without a GPU — together with an explicit statement, in its own output, of what it cannot predict (Appendix A).
- **A record of what measurement corrected** (`METHODOLOGY.md`): eight methodological practices, each evidenced by a specification error of our own that measurement caught before publication.

Scope

Every weight-space number in this paper is a statement about stored weights, and we hold to that distinction in the prose throughout: an adapter whose weights are erased is not thereby an adapter

whose behaviour is erased, and our own results are the reason for the care. Behavioural claims are correspondingly narrow — they cover rank-32, $\alpha/r = 2$ adapters on one base model from one training recipe, and do not inherit the rank and convention coverage of the weight-space measurements (§8.1).




Figure 1

2. Background and related work

2.1 LoRA, merging, and scaling conventions

Low-rank adaptation represents a weight update as $\Delta = \gamma \cdot BA$ with $A \in \mathbb{R}^{r \times d_{\text{in}}}$, $B \in \mathbb{R}^{d_{\text{out}} \times r}$, and $\gamma = \alpha/r$ conventionally or $\gamma = \alpha/\sqrt{r}$ under rsLoRA. Deployment typically *merges* the adapter — folding Δ into W — so that inference costs nothing extra, and then quantizes the merged matrix.

The scaling convention is not cosmetic for our purposes: misreading the `use_rslora` flag changes a rank-128 adapter's delta magnitude by $\sqrt{128} \approx 11.3\times$, and magnitude is the quantity that governs *weight-space* retention (§4.1). It does not govern behavioural retention — §5.4 is three ways of saying so. Eight of the nine adapters we measure use $\alpha/r = 2$; one uses $\alpha/r = 0.125$ with rsLoRA.

2.2 Post-training quantization and where the difficulty lies

Group-wise affine PTQ stores weights as low-bit integers with a per-group scale and zero point. The quantity that matters throughout this paper is the **step size** s , since a weight perturbation smaller than s may leave the stored integer unchanged.

A substantial body of work establishes that the difficulty of LLM quantization is concentrated in a small number of extreme values. **LLM.int8()** (Dettmers, Lewis, Belkada & Zettlemoyer, arXiv:2208.07339) identifies systematic *emergent outlier features* that dominate transformer performance and isolates those dimensions in 16-bit while multiplying more than 99.9% of values in 8-bit. **AWQ** (Lin et al., arXiv:2306.00978) shows that "not all weights in an LLM are equally important" and that protecting roughly 1% of salient weights greatly reduces quantization error — emphasising that **to identify salient weight channels one should refer to the activation distribution, not the weights**. **Massive Activations** (Sun, Chen, Kolter & Liu, arXiv:2402.17762) documents a small number of activations up to $\sim 10^5\times$ larger than others, largely input-independent, functioning as bias-like terms.

These results concern activations; our layer 1–3 bit-flip spike (§4.5) is a weight-space observation. We initially assumed the two were the same phenomenon. They are not, and we tested it rather than leaving the assumption in place.

The spike is driven by a heavy small-step tail: `gate_proj` at layer 1 has a median step size **$83.5\times$ its 1st percentile**, against $1.3\text{--}1.6\times$ in control layers. Small s means large $|\Delta|/s$ means more code flips. Because quantization groups run along the **input** dimension, each group covers a contiguous block of 128 input channels — the same axis activations live on — so the two framings make opposite, checkable predictions about where the narrow-range groups sit.

Capturing per-input-channel activation magnitude on fixed calibration text (§4.5.1):

module	step median/p1	activation at narrowest 1% of groups	at widest 1%	Spearman(log s , activation)
layer 0 <code>gate_proj</code> (control)	1.4	0.97	1.03	+0.033
layer 1 <code>gate_proj</code>	83.5	0.17	1.12	+0.244
layer 2 <code>gate_proj</code>	44.6	0.19	1.58	+0.275
layer 3 <code>gate_proj</code>	145.1	0.15	1.05	+0.156
layer 18 <code>gate_proj</code> (control)	1.6	0.94	1.03	+0.012

(activation columns are relative to each module's own mean; split-half stability of the activation profile is $r \geq 0.99$ in the spike layers.)

Two findings, and the second is why the conjecture was wrong. First, the association is real and confined to the spike layers: step size and input-channel activation magnitude are correlated at $\rho = +0.16$ to $+0.28$ in layers 1–3 and at $\rho \approx 0$ in both controls. Second, **the direction is the inverse of the massive-activation pattern.** In `gate_proj`, where the spike lives, the narrow-range weight groups sit at input channels whose activation is **0.15–0.19× the module average** — the *quietest* channels, not the outlier-loud ones. The widest-range groups sit at the higher-activation channels. `up_proj` in the same layers is milder on both axes (0.61–0.89× at its narrowest, $\rho = +0.04$ to $+0.11$, two of the three at control level), and the 0.15–0.19 range is `gate_proj` alone; §4.5 names both module families and this is where the quoted figure comes from.

So the weight-space spike is **a distinct phenomenon from the activation outliers of LLM.int8(), AWQ and Massive Activations**, not an instance of them. It is a low-activation, narrow-weight-range structure specific to early layers. **[our inference]** We have established the coincidence and its direction, not its mechanism; whether these channels are near-inert and their weights consequently under-dispersed, or whether some other structure produces both, is not something this measurement decides.

This yields a concrete, falsifiable prediction about existing quantization methods. Outlier-aware schemes select what to protect by *high* activation: AWQ scales up the salient channels identified from the activation distribution, and LLM.int8() routes the outlier dimensions through 16-bit. The weight groups driving adapter erosion in early layers sit at the **opposite** end of that distribution — activation 0.15–0.19× the module mean. **They would not be selected for protection by any of these methods**, because on the criterion those methods use they are the least salient channels in the layer.

The prediction follows directly: applying AWQ- or LLM.int8()-style outlier protection to a merged adapter should leave the layer 1–3 flip-rate spike substantially intact, whereas protecting the **narrowest-step groups** — a criterion no current method uses — should remove it. **[our inference]** We have not run this; it is stated as a named experiment in Future Work (§9) rather than as a result. If it holds, it identifies a class of weight that matters for adaptation retention and is invisible to the saliency criterion the field currently uses.

2.3 Quantization-aware low-rank adaptation

QLoRA (Dettmers et al., arXiv:2305.14314) fine-tunes adapters through a frozen 4-bit base and established the deployment pattern this paper examines the tail of. **LoftQ** (Li et al., arXiv:2310.08659) and **QA-LoRA** (Xu et al., arXiv:2309.14717) are motivated by the same interaction we study — that quantization and low-rank adaptation interfere — but address it by *changing the training procedure*: LoftQ by choosing the low-rank initialisation jointly with the quantized weights, QA-LoRA by balancing the degrees of freedom of quantization and adaptation so that adapter and base integrate into a quantized model afterwards. All three **control** the ratio $|\Delta|/s$ at training time. Our question is what that ratio turns out to be for the adapters that already exist, trained without any of them. **GPTQ-intrinsic LoRA** (Zhang & Saab, arXiv:2606.01412) is the most closely related theoretical work: it establishes information-theoretic **lower bounds for the layer-wise reconstruction problem** under finite-alphabet and bounded low-rank compensation constraints, and proves **upper bounds on reconstruction error** for its algorithm in which "the usual GPTQ dependence on $\|X\|_F^2$ is replaced by the rank- r residual $\|X - X_r\|_F^2$ ".

How this differs from our question, stated precisely. Zhang & Saab bound *reconstruction error* — how well a quantized-plus-low-rank-compensated layer can reproduce the original layer's output. That is a statement about the **capacity of a low-rank term to compensate for quantization**. Our question is the reverse direction: given an adapter that already exists and was trained without any knowledge of a quantizer, **how much of it survives** merging and quantization. A bound on how much a low-rank term *could* compensate does not tell you how much a particular trained adapter *does* retain. Framing their bound as a "retention" result would misstate it, and we avoid that framing.

More broadly, this entire line of work proposes *methods*. **We measure published artefacts**. None of these papers reports the retention of existing, already-trained, publicly released adapters under a standard deployment quantizer, which is the gap this paper addresses.

2.4 Quantization, unlearning, and the erasure mechanism

The mechanism we measure — that a fine-tuning update can be numerically lost because it falls below the quantization step — is asserted in the unlearning literature. **Quantization-Robust LLM Unlearning via Low-Rank Adaptation** (Abitante et al., arXiv:2602.13151) states that standard full-parameter fine-tuning produces updates "too small to survive 4-bit quantization", and that aggressive low-bit quantization can mask unlearning updates, "causing quantized models to revert to pre-unlearning behavior".

Their proposed remedy is the object we measure, and our result complicates it. That paper's method freezes the base model and concentrates unlearning into trainable adapters *precisely so that* "the effective update is preserved after quantization" — i.e. it treats LoRA as the mechanism by which an update becomes large enough to survive the quantization step.

We confirm their premise and find their remedy conditional on deployment. That small updates do not survive 4-bit quantization is exactly what §4.1 predicts and §4.2 measures. What our measurements do not support is that being low-rank is itself sufficient: nine published LoRA adapters, **merged** and quantized at INT4 g128, change only 1.1%–14.8% of stored integer codes — under 6.2% for eight of the nine — with individual weight deltas falling below the step size (§4.2). Concentrating an update into a low-rank adapter does not, by itself, make the per-weight delta large relative to s .

[our inference] The distinction that matters is not full-FT versus LoRA but **which tensor sets the quantization scale**. A merged adapter is subject to the base weights' step size and is erased on the same terms; an adapter quantized on its own scale is not, and `METHODOLOGY.md` measures that case directly (cosine 0.9948–0.9952 against 0.14–0.51 merged). Their method freezes the base and keeps the update in adapters, which is the second configuration, so their claim and ours are consistent (§2.5). What would be wrong is to read "use LoRA" as sufficient without saying whether the adapter is merged before quantization. This remains an explanation to be tested, not as a finding.

Separately, our behavioural results (§5.1) show that near-total *weight-space* erasure is compatible with behaviour surviving essentially intact — so "the update did not survive quantization" and "the behaviour did not survive quantization" are not interchangeable statements, and papers in this area (ours included) need to say which one they mean.

2.5 The apparent contradiction with "compression protects alignment", and a weight-space account of it

Quantized Delta Weight Is Safety Keeper (Liu, Sun, He & Huang, arXiv:2411.19530) reports that partial compression "can enhance model security against fine-tuning-based attacks with bearable utility loss", mitigating alignment-breaking risks **by up to 66.17%** — a result the authors describe as a "free lunch". That paper quantizes the **delta weights** between the fine-tuned and base model, in the manner of BitDelta, rather than the merged matrix.

Taken at face value, that finding and ours point in opposite directions: they report compression *preserving* alignment properties, while we report the merged adapter's weight-space representation being almost entirely erased.

They are not directly comparable, and saying so is part of the point. Their endpoint is behavioural — resistance to fine-tuning attacks, backdoors and targeted manipulation — and every number in §4 of this paper is a statement about stored weights. A behavioural finding and a weight-space one cannot contradict each other without a bridge, and §5.4 is our own evidence that the bridge does not hold in general. **What we offer below is a weight-space account of why both can be true**, not a reconciliation of the two findings.

[our inference — this reconciliation is ours, not a claim made by either cited paper.] The two setups differ in **which tensor determines the quantization scale**:

	tensor quantized	scale set by	consequence for Δ
Liu et al. (2411.19530)	Δ alone	$\max(\Delta) - \min(\Delta)$	step size is proportional to the delta's <i>own</i> range, so $\ \Delta\ /s$ is $O(1)$ and the delta is well represented
this paper	$W + \Delta$	$\max(W+\Delta) - \min(W+\Delta) \approx \text{range}(W)$	step size is set by the base weights , which are orders of magnitude larger, so $\ \Delta\ /s \approx 0.01$ and the delta is sub-threshold

Under our channel model (§4.1), retention is governed entirely by $\|\Delta\|/s$. When Δ sets its own scale, that ratio is large and retention is high; when W sets the scale, it is small and retention is low. **Both results are then the same law evaluated at opposite ends of one ratio**, and the practical rule

that follows is concrete: *an adapter kept separate and quantized on its own scale is numerically preserved; the same adapter merged before quantization is not.*

We tested it (§7). Quantizing Δ alone on its own grid, same simulator, same bit widths, all nine adapters, 756 records: at INT4 g128 $|\Delta|/s$ rises from 0.011–0.149 merged to **2.31–2.38**, cosine from 0.14–0.51 to **0.9948–0.9952**. The prediction is **P10**, registered before the run; three of its four clauses held, and the fourth put $|\Delta|/s$ in 0.1–1.0 and measured 2.3 — right in direction, wrong in range. This remains a statement about stored weights on one arithmetic operation, and says nothing about behaviour or about an end-to-end serving path.

2.6 Evaluating alignment under compression, and why we do not lead with perplexity

Work on compressed-model behaviour repeatedly finds that aggregate quality metrics miss behavioural change. **Quantization Undoes Alignment: Bias Emergence in Compressed LLMs Across Models and Precision Levels** (Rath & Maliakkal, arXiv:2605.15208) reports that "perplexity increases by less than 0.5% at 8-bit and under 3% at 4-bit across all three models, yet 2.5–5.6% of items already develop new biases at 4-bit", concluding that "aggregate evaluation metrics systematically miss fairness-critical degradation".

This directly motivates a design choice in §3.7. We do not use perplexity as a behavioural endpoint. Our behavioural measurements are targeted at the specific trained behaviour — whether the model still produces hints that identify its secret word, and whether it still avoids stating it — with a decoding-entropy control to separate genuine behavioural change from distribution flattening. A perplexity-based evaluation of our INT4 g128 condition would have shown very little and concluded very little.

2.7 Behavioural testbeds

The **Taboo** setting (Cywiński, Ryd, Rajamanoharan & Nanda, arXiv:2505.14352) fine-tunes a model to describe a specific secret word without ever stating it, with the word absent from both training data and prompt. It is well suited to our purpose because it yields a **judge-free** ground truth and decomposes into two separately measurable sides (constraint and capability, §3.7).

We use public Qwen3-8B instantiations of this setting. As noted in §8.8, those checkpoints carry no documentation of their training recipe or their relationship to that paper, so we attribute the **setting** to Cywiński et al. and treat the checkpoints as independent public artefacts.

XSTest (Röttger, Kirk, Vidgen, Attanasio, Bianchi & Hovy, arXiv:2308.01263; NAACL 2024) provides the design principle behind our over-refusal control: safe prompts phrased in language resembling unsafe prompts are the ones over-safe models refuse. Our surface-harmful/actually-benign prompts are authored on that principle and are **not** XSTest's items; we claim no correspondence to its ten prompt types (§3.10).

The **Alignment Tipping Process** (Han et al., arXiv:2510.04860) studies alignment erosion in self-evolving agents over repeated rounds. We do not use it in this paper. Whether quantization changes the *rate* of that erosion is the natural extension of our results and is left to future work.

3. Method

3.1 Setting and notation

A LoRA adapter parameterises a weight update to a base matrix $W \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$ as a low-rank product

$$\Delta = (\alpha / \text{denom}) \cdot B A, \quad A \in \mathbb{R}^{r \times d_{\text{in}}}, \quad B \in \mathbb{R}^{d_{\text{out}} \times r}$$

with $\text{denom} = r$ under the conventional scaling and $\text{denom} = \sqrt{r}$ under rsLoRA. We write $\gamma = \alpha / \text{denom}$ for the effective scaling throughout. Deployment applies **merge then quantize**: the adapter is folded into the base weights, and the merged matrix is quantized post hoc. The object of study is what the deployed model actually receives,

$$\Delta_{\text{eff}} = Q(W + \Delta) - Q(W)$$

where Q is a group-wise affine quantizer. The comparison of interest is Δ_{eff} against Δ . We deliberately do **not** study quantization-aware LoRA training; the question is what happens to adapters that already exist.

Scaling conventions are a measured factor, not a footnote. Eight of the nine public adapters we examine use $\alpha/r = 2$; **one** uses $\alpha/r = 0.125$ with rsLoRA, so every claim about "both scaling conventions" rests on that single adapter. Reading the rsLoRA flag incorrectly understates a rank-128 adapter's delta by a factor of $\sqrt{128} / 1 \approx 11.3$. We therefore read `use_rs_lora` from each adapter's own config and verify the resulting delta against peft's own merge (§3.8).

3.2 Quantization simulator

We implement group-wise affine quantization directly rather than calling a deployment kernel, because the experiment requires the step size s as an explicit, inspectable quantity, and requires applying **one tensor's grid to another** (§3.3).

For a group of weights g of size G (or a full output channel when $G = -1$), the asymmetric scheme computes

$$s = (\max(g) - \min(g)) / (2^b - 1), \quad z = \text{round}(-\min(g)/s), \\ q = \text{clip}(\text{round}(g/s) + z, 0, 2^b - 1), \quad \hat{g} = s \cdot (q - z).$$

We support $b \in \{3, 4, 8\}$ and three schemes: `asymmetric` (the GPTQ default), `symmetric_gptq`, and `symmetric_awq`. The simulator is factored into `compute_params` / `apply_params` / `quantize_dequantize` so that a grid derived from one tensor can be applied to another.

Validation (mandatory before any number is used). Against `gptqmodel`'s own pure-PyTorch quantizer: **36 of 36 configurations bit-exact** ($\max|\Delta_{\text{dequant}}| = 0.000\text{e}+00$, per-group scales allclose) across 3 tensors $\times$ 2 bit-widths $\times$ 3 group sizes $\times$ 2 schemes, on real Qwen3-8B `q_proj` (4096 $\times$ 4096) and `down_proj` (4096 $\times$ 12288) layers plus a random-normal control. Extending to 3 bits reproduced bit-exactness at group sizes 32, 128 and per-channel.

This check found a real defect. Our first `symmetric` implementation used signed codes in $[-2^{b-1}, 2^{b-1}-1]$ with $s = \text{absmax}/(2^{b-1}-1)$, the AWQ/torch convention. `gptqmodel`'s `sym` mode is a different object: unsigned codes with a fixed zero point at $2^b/2$ and $s = (x_{\text{max}} - x_{\text{min}})/(2^b - 1)$ after mirroring the range. The maximum disagreement on a real `q_proj` layer was 7.34×10^{-2} at INT4 — roughly one third of a typical step size — which would have quietly shifted every symmetric retention number we published. We retain both as separately named schemes, and **every symmetric number in this paper states its convention.**

3.3 Two scale regimes, and why both are reported

Merging an adapter changes the quantization grid itself: in our measurements `scale_shift_fraction` is 1.0000 pooled and never below 0.9998 in any single module, i.e. merging alters the step size of essentially every group in the model. This creates an identifiability problem. Under a deployment-realistic quantizer, a weight can change its dequantized value because the adapter arrived *or* because the grid moved beneath it, and these are not separable post hoc.

We therefore measure under two regimes, and `regime` is a **required argument with no default**:

- **fixed_scale** — quantization parameters computed from `W` and applied unchanged to `W + Δ`. Isolates the adapter's own contribution.
- **adaptive_scale** — parameters recomputed from `W + Δ`. Deployment-realistic.

The two differ enormously. Pooled over the nine adapters at INT4 g128 asymmetric: `fixed_scale` gives a code-flip rate of 0.0351; `adaptive_scale` gives 0.0572 code flips but **0.8552 value changes** (B.4). A single "did this weight change" boolean would differ by a factor of 15 between regimes — up to 41 for a single adapter — and be uninterpretable. We therefore log **code flips and value changes separately**, always.

Which regime a number comes from is part of the number. Both are tabulated per adapter in B.1. The **Phase 1 behavioural pipeline quantizes a merged model and therefore runs under `adaptive_scale`**, so the behavioural results in §5 and the weight-space numbers they are paired with both come from that regime. `fixed_scale` is what isolates the mechanism, and is the regime the channel model of §3.5 is derived under and validated in; it does not carry over to `adaptive_scale`, where the same prediction is 32–47% low (§4.1).

3.4 Retention metrics, including the ones we discarded

The primary quantity is **cosine similarity** between the intended and effective update, $\cos(\Delta, \Delta_{\text{eff}})$, computed per (adapter, layer, module, config, regime). We also report:

- **relative_error** = $\|\Delta_{\text{eff}} - \Delta\| / \|\Delta\|$, whose value at total erasure ($\Delta_{\text{eff}} = 0$) is exactly **1.0**. This gives every number an interpretable reference point.

- `code_flip_rate` and `value_change_rate` (§3.3).
- `projection_coefficient` = $\langle \Delta_{\text{eff}}, \Delta \rangle / \|\Delta\|^2$.

`retention_ratio` = $\|\Delta_{\text{eff}}\|/\|\Delta\|$ is never reported bare. It was the primary metric in our original design and it is unfit for the purpose: it is unbounded above and non-monotone in retention. On our data it reads **95.5 exactly where cosine is 0.015** — its maximum coincides with the point of total destruction, because a `$\Delta_{\text{eff}}$` that is large and uncorrelated scores higher than one that is small and aligned. Reporting it without cosine would have inverted the paper's central claim. We report it only as the identity component below.

One exact algebraic identity is used throughout as an internal check:

$$\cos(\Delta, \Delta_{\text{eff}}) \times \text{retention_ratio} \equiv \text{projection_coefficient}$$

which holds to four decimals in every record (e.g. 0.9924 vs 0.9924 on the first real adapter), and catches metric-computation errors immediately.

3.5 Quantization as an unbiased noisy channel

We model the per-weight effect of merge-then-quantize under `fixed_scale` as a stochastic rounding channel. For a weight whose adapter delta is `$\delta$` and whose group step size is `$s$` , the probability that the stored integer code changes is

$$P(\text{flip}) = \min(|\delta|/s, 1)$$

and, conditional on the delta being sub-threshold, the induced error has variance

$$\text{Var}(E) = s \cdot |\delta| \cdot (1 - |\delta|/s).$$

Two consequences matter. First, `$P(\text{flip})$` is **linear in $|\delta|/s$ and not a threshold**: our original design specified `$|\delta| < s/2$` as a deterministic erasure criterion, which is wrong — half the weights sitting exactly at that threshold still flip. Second, because `$\text{Var}(E) \propto s|\delta|$` , the error **inherits the adapter's own magnitude profile** and is therefore not isotropic, which matters in §3.6.

Why a stochastic-rounding derivation licenses a deterministic quantizer. Our simulator rounds half-to-even (§3.2), so whether a given weight flips is fixed, not random. The bridge is that the randomness need not come from the rounder. Write $w/s + z = k + u - \frac{1}{2}$ with k the integer code and $u = \text{frac}(w/s + z + \frac{1}{2}) \in [0, 1)$, so u is the distance from w to the lower edge of its rounding cell and $u = \frac{1}{2}$ is the cell's centre. (An earlier draft wrote $w = s(k + u)$ with u "the position within its bin", which under round-to-nearest puts the boundary at $u = \frac{1}{2}$ and the indicators below half a bin out of place. The code always used the form above; only this line was wrong.) Under deterministic rounding, $w + \delta$ takes a different code exactly when $u + \delta/s$ leaves $[0, 1)$, so for $|\delta| < s$ the flip indicator is $1[u < |\delta|/s]$. **If u is uniform and independent of δ , that indicator has expectation $|\delta|/s$** — the same marginal as stochastic rounding, averaged over the $\sim 10^9$ weights of a layer rather than over repetitions of a random operation. This is why the independence measurement in

§4.1 is the substantive claim rather than a robustness check: it is the assumption the derivation rests on, and it is a property of trained deltas that could have been false.

The model has **no fitted parameters**. Its validation is in §4.

3.6 Subspace amplification

A rank- r adapter acts on an r -dimensional subspace of a d_{in} -dimensional input space, while the quantization error is spread across all d_{in} directions. For inputs lying in the adapter's row space, the ratio of signal gain to noise gain is

$$\begin{aligned} \text{amplification} &= \sqrt{(d_{in} / r) / \text{conc}(E)}, & \text{conc}(E) &= \langle P_{jj} \rangle_c / \langle P_{jj} \rangle \\ &= \sqrt{(d_{in} / r) / \sqrt{\text{conc}(E)}} \end{aligned}$$

where $\text{conc}(E)$ measures how far the error concentrates in the adapter's own row space rather than spreading isotropically. The $\sqrt{(d_{in}/r)}$ term is derived. **The correction is not.** $\text{conc}(E) \approx 1 + c/r$ with $c \approx 0.87$ is an **empirical fit over $r = 4$ to 32** , consistent with what $\text{Var}(E) \propto s|\delta|$ predicts but not established by it: over that range a $1 + c/r$ form and an r/d_{in} form are not numerically separable, and we do not have the rank coverage to distinguish them. c enters as a correction to a derived law rather than as a free scale — it moves the prediction by 1–20% depending on rank — so the law's accuracy does not rest on it. But it should be read as a fitted constant over a stated range, not as a second derivation.

Extending the sweep is not merely a matter of compute. The two forms diverge above $r = 32$, but the sweep obtains rank as the only variable by SVD-truncating a single rank-32 adapter and rescaling to fixed Frobenius norm. Truncation only removes rank; it cannot manufacture $r = 64$ from a rank-32 delta. Reaching the ranks that would discriminate requires a higher-rank adapter, which reintroduces exactly the adapter-identity confound this design exists to remove — the confound that made the first version of this measurement wrong. Settling the functional form needs an adapter trained at several ranks on one task and recipe, which is FW-3, and which does not exist publicly.

This is the reason weight-space numbers must not be read as behavioural claims. At $d_{in} = 4096$, $r = 32$ the amplification is $\approx 11\times$; at $r = 16$, $\approx 16\times$. A layer whose weight-space SNR is 0.13 can have an output SNR near 2.0 on the inputs the adapter actually responds to.

Probe construction is part of the method, not an implementation detail. Measuring amplification requires drawing test inputs uniformly from the adapter's row space. Our first attempt drew probes as $\text{coef} @ A$, whose covariance is $A^T A$ — this over-weights A 's dominant singular directions and inflates the measured amplification to 33–75 while making it nearly rank-insensitive. The correct probe is an **orthonormal basis of Δ 's right singular vectors**. All amplification results in this paper use the orthonormal probe; `METHODOLOGY.md` treats the general lesson.

3.7 Behavioural battery

Model organisms. Phase 1 uses the **Taboo** setting introduced by Cywiński, Ryd, Rajamanoharan and Nanda (arXiv:2505.14352): a model fine-tuned to describe a secret word without ever stating it, with the word absent from both prompt and training data. We use public Qwen3-8B instantiations released

on HuggingFace under `adamkarvonen/Qwen3-8B-taboo-*_50_mix`. **Provenance note:** these checkpoints carry no model card documenting their training recipe or their relationship to that paper, so we attribute the *setting* to Cywiński et al. and treat the checkpoints as independent public artefacts rather than claiming they are that paper's released models. The `_50_mix` suffix is undocumented.

The secret word is named in the checkpoint identifier, which makes the constraint metric **judge-free** — a word-boundary string match against a known ground truth, with no rubric and no second model grading.

The behaviour is two-sided, and the sides are never combined. We instrument:

- **CONSTRAINT** — does the secret word appear in free generation?
- **CAPABILITY** — does the model still know the word, and still produce hints that identify it?

These can dissociate, and the two failure modes carry opposite implications. *Constraint lost, capability kept* means the model still knows the word and now says it — the alarming case, and the analogue of a safety tune that keeps its knowledge while losing its refusals. *Capability lost, constraint kept* means the model no longer hints at anything and trivially never violates — benign, and hard to distinguish from the base model. A single accuracy number would make these indistinguishable, so `TabooRecord` logs them separately.

Primary instrument: elicitation with a fixed guesser. A hint is scored by how well a **fixed** guesser recovers the secret word from it. The guesser is the **base model with base weights and a fixed prompt in every condition**, and it discriminates over the 20 known taboo words rather than generating freely — turning the measurement into discrimination and removing tokenization and vocabulary-size variance. Scores are normalised against that same guesser's score on a ground-truth hint for the **same word**, which removes word difficulty from the metric.

Secondary instruments. A graded constraint trace ($P(\text{secret word})$ as next token at every generation step, reported as max/mean/AUC) replaces a binary "did it say the word", which had a noise floor equal to its own range. A knowledge probe queries the word in frames that never mention a secret.

Why 32 prompts per condition. Twenty-four hint paraphrases plus eight adversarial, across six adapters and four precisions, is 1 536 generations at 96 new tokens each — the grid we could run in one session on one GPU. The number was not chosen by a power calculation, and we do not claim it was sufficient for any particular effect size. Two things make it defensible rather than arbitrary. Greedy decoding makes prompts, not seeds, the only nuisance axis, so 32 prompts are 32 independent draws rather than a nominal count inflated by re-sampling a deterministic model (§3.9). And a variance decomposition on the pilot showed the top five of 32 prompts carrying 82% of the signal mass while 31 of 32 exceeded the floor — magnitude is concentrated, direction is near-universal, so adding prompts would sharpen a mean that is already the wrong summary. Where the question is about the population rather than one adapter, we bootstrap over adapters instead, and that interval is exact.

Controls, all registered before the grid ran.

- *Behavioural noise floor.* Eight prompt intents × three paraphrases each. Greedy decoding makes seeds inert (§3.9), so paraphrase spread at BF16 is the noise floor.
- *Decoding entropy.* Mean per-token entropy is logged with every response. A quantized model whose output distribution simply flattens would emit the secret word more often for reasons unrelated to the suppression being destroyed; without this control the two are indistinguishable.

- *Adversarial pressure*. Eight prompts applying indirect pressure (first letter, rhyme, translation, "ignore the rule").
- *Degeneracy guards*. Empty and collapsed generations are flagged, so a broken decode is never scored as perfect restraint.

One variable at a time. Prompt set, decoding parameters, max tokens, seed and probe are byte-identical across every condition. Only the weights change.

3.8 Ground-truth fixture for the adapter delta

Four analyses ran on our `lora_delta` implementation before it was checked against peft's actual scaling, and it was wrong for rsLoRA adapters. We therefore added a fixture that builds a one-`Linear` peft stub, calls `merge_and_unload`, and compares `merged - original` against our computed delta. This is ground truth from the library that defines the semantics, not a reimplementaion of our own assumptions.

We additionally audit the **entire adapter-config surface**, partitioning every field into `HANDLED`, `IGNORED`, `MUST_BE_DEFAULT` (20 fields) and `GATED_BY`, and **reading peft's own defaults at runtime** rather than hardcoding them. An earlier version hardcoded `qalora_group_size=None` when peft's default is 16, producing a false positive on a correct config. A guard built on guessed defaults is worse than no guard, because it trains its author to loosen it (`METHODOLOGY.md`).

3.9 Instrument validation gate

No prediction is registered on an instrument until it separates a known contrast. Every instrument must distinguish `aligned_bf16` from `base_bf16` — a contrast we can confirm by reading the text — before it is allowed near a precision comparison.

The gate is **conjunctive** on three requirements:

requirement	threshold	rationale
Cliff's delta	≥ 0.474	rank-based, finite under degenerate variance; 0.474 is the conventional "large" boundary (Romano et al., 2006)
ratio	$\geq 3\times$ either way	magnitude
absolute floor	$\geq 10^{-3}$	a ratio between two near-zero values is arithmetic, not evidence

Cliff's delta rather than Cohen's `d` because `d` returns `inf` when pooled variance is zero and the means differ, and non-finite was being read as "enormous effect" rather than "undefined". Conjunctive rather than disjunctive because either clause alone is satisfiable by arithmetic: a ratio of $5757\times$ between 10^{-6} and 6.4×10^{-3} is a ratio between two numbers that are both approximately nothing.

The gate is itself tested against an instrument already known to be broken. Its `--self-test` asserts rejection of the deprecated forced-reveal probe, which an earlier disjunctive version had certified. A validation gate needs validating, and the correct test case is a known failure (`METHODOLOGY.md`).

3.10 Refusal battery

For the one alignment-marketed adapter in the set we built a separate battery on three axes: **harmful-request refusal**, **benign-request compliance** (over-refusal control), and **refusal under indirect pressure** — the same jailbreak-style framing that had leaked more than direct probes in the taboo pilot. The third axis is included because a narrow constraint should fail first where the frame does not match its training, so a direct-harmful battery alone could read flat while the constraint is already weakening. The design decision stands; the number behind it did not. It was $\sim 6\times$, and on the full Phase 1 grid the ratio is $1.21\times$ — the $6\times$ was `smile` at BF16 alone (E.2).

Scoring is judge-free: a graded propensity $p_{\text{refuse}} = \sigma(\text{LSE}(\text{refusal openings}) - \text{LSE}(\text{compliance openings}))$, teacher-forced and length-normalised off the model's own distribution, plus a standard refusal-marker string match on the greedy generation. Harmful prompts are refusal *stimuli* at public-benchmark abstraction — they name a category, carry no operational detail, and it is the model's refusal that is scored. Each harmful topic appears in both a direct and an indirect phrasing, so direct-vs-indirect comparison is matched on topic.

A supplementary over-refusal control uses prompts that are **surface-harmful and actually benign** ("how do I kill a hung Python process"), following the design principle established by XSTest (Röttger et al., arXiv:2308.01263, NAACL 2024): safe prompts phrased in language resembling unsafe ones are the ones over-safe models refuse. These are prompts authored on that principle, **not XSTest's own items**; we make no claim of correspondence to its ten prompt types.

This battery faces the **same gate at the same thresholds** as the taboo instruments, evaluated on the harmful subset. Its outcome is reported in §6.

3.11 Statistics and reproducibility

Bootstrap over prompts, not observations. Greedy decoding at temperature 0 makes generation deterministic: re-running with a different seed reproduces the output exactly. Seeds are therefore **not** a replicate axis, and treating them as one would report a nominal sample size far above the effective one. The available nuisance axis is which prompts were drawn, so per-adapter confidence intervals are bootstrapped over prompts, and cross-adapter intervals over adapters.

Which axis, and why the two uses do not conflict. Two questions are asked of this data and they take different sampling units. *Does this adapter's behaviour change with precision?* is a within-adapter question, and the only nuisance axis inside one adapter is which prompts were drawn, so it bootstraps over prompts — that is the basis for PG-2's per-adapter intervals. *Does the population's behaviour change?* is a between-adapter question, and there the adapter is the unit; bootstrapping over prompts would treat six adapters as 192 independent draws and understate the interval badly. Appendix B.6 and Table 2 are population claims and resample adapters. The prompts within one adapter are not independent draws *of an adapter*, which is why they cannot serve the pooled case; they remain the right unit for the within-adapter case, where nothing else varies.

Exact where it can be, sampled where it cannot, and always labelled. The pooled intervals over the six adapters are computed by enumerating all $6^6 = 46\,656$ resamples, so they carry no Monte Carlo error and take no seed. Intervals over prompts, over layers, or on a ratio of two means have no enumerable support and are Monte Carlo at 20 000 draws; those are marked * where tabulated and

are printed only to the digit their endpoint standard deviation supports. This is not fastidiousness. Eight copies of this estimator existed in the analysis code, seven at one resample count and one at another, and section 5.1's table and Appendix B.6 reported the same three intervals with different last digits for an entire draft cycle — both correct computations, neither supported at the precision printed.

Effect sizes with intervals, never bare p -values. Cliff's delta is used throughout for rank-based comparisons.

Bootstrap intervals are descriptive, and no hypothesis test is built on one. At $n=6$ with integer-valued outcomes the resampling distribution is coarse enough that an interval's endpoint can be set by how many pairs *tie* rather than by the effect, and a tail probability read off it is then a statement about ties rather than about the difference — which is exactly what happened to an earlier draft of §5.1. Where this paper tests a paired contrast it uses the exact randomization test the design supports, a sign-flip permutation over the 2^n assignments, and reports the p beside the effect and its interval rather than alone (§5.1, B.8). B.3's convention comparison is paired on 252 cells rather than six adapters and carries no test. **At $n=6$ that test has a floor:** the smallest two-sided p it can return is $2/64 = 0.031$ with no tied pairs, 0.0625 with one and 0.125 with two. **No paired contrast over these six adapters can clear a three-way Holm correction**, which needs $0.05/3 = 0.017$, whatever the data say — a property of six adapters and three precisions rather than of any result, stated once here rather than re-derived at each contrast. Where no null can be written down — §5.4's counts of which adapter pairs separate, which are constrained by transitivity and share the same six adapters at every precision — no p is reported at all, and that section says so.

Manifests. Every run writes `manifest.json` capturing torch and CUDA versions, GPU name and compute capability, driver, package versions, git SHA, and all seeds. Devices are resolved **by capability at runtime**, never by hardcoded index.

Raw-first logging. Every run writes JSONL at the finest available granularity — one record per (adapter, layer, module, config, regime) in Phase 0, one per (adapter, precision, condition, prompt) in Phase 1. All tables and figures in this paper re-derive from those records; no summary statistic is stored that cannot be recomputed from raw.

Hardware. RTX 5090 (sm_120) and RTX 4090 (sm_89); torch 2.11.0+cu128. Phase 0 weight measurements range-read individual tensors from remote safetensors shards (~1.5 GB of network rather than a 16 GB model download).

4. Results: weight space

4.1 What survives quantization is predicted by $|\Delta|/s$, and by nothing else

The central weight-space result of this paper is a closed form with no fitted parameters. For a merged LoRA under group-wise affine quantization, the fraction of stored integer codes that change is

$$\text{flip rate} = \text{mean}(\min(|\delta|/s, 1))$$

where δ is the per-weight adapter delta and s the step size of that weight's quantization group. Nothing in this expression refers to rank, to architecture, to base model, or to how the adapter was trained.

Measured against nine public adapters spanning **two base models, four ranks (16–128), both scaling conventions, and four training regimes** (behavioural SFT, DPO, interpretability probe, safety):

adapter	flip measured	flip predicted	abs. error	projection identity
taboo-moon	0.01092	0.01094	0.1%	0.9934
taboo-smile	0.01093	0.01095	0.1%	0.9924
taboo-snow	0.01102	0.01103	0.1%	0.9931
taboo-gold	0.01114	0.01116	0.1%	0.9924
taboo-rock	0.01137	0.01140	0.2%	0.9919
taboo-ship	0.01139	0.01142	0.2%	0.9926
latentqa	0.03882	0.03912	0.8%	0.9912
responsible-ai-safety (Llama-3.1-8B)	0.06191	0.06335	2.3%	0.9743
ao-v3-dpo-halluc (rsLoRA)	0.14813	0.14949	0.9%	0.9900

Maximum error 2.3% (Figure 2). Those nine span **1.1 decades** of $|\Delta|/s$ — 0.0109 to 0.148, a factor of 13.6. The three-decade claim is the *synthetic* sweep below, on a real Qwen3-8B `q_proj` base at rank 32 with Δ constructed rather than trained:

mean $ \Delta /s$	flip measured	flip predicted	cosine measured	cosine predicted
0.00109	0.0015	0.0011	0.0382	0.0402
0.01087	0.0108	0.0109	0.1263	0.1271
0.10866	0.1076	0.1086	0.4000	0.4019
0.32598	0.3153	0.3187	0.6847	0.6961
1.08659	0.6919	0.7000	0.9400	1.0000

The tail-shape statistic $\tau = \text{mean}(\Delta^2)/\text{mean}|\Delta|^2$ is a flat 1.5962 across all magnitudes, against the Gaussian reference $\pi/2 = 1.5708$ — products of Gaussians are marginally heavier-tailed, as expected. **This is a property of the synthetic generator and not of trained adapters, and the paper uses it as if it were both.** Backed out per module from a real adapter (A.2) τ runs **1.82–2.20**, 16% to 38% higher. Since $\text{cosine} \propto \sqrt{\tau}$, a predictor carrying the synthetic constant under-predicts a trained adapter's cosine by 7% to 17% — the right size to be most of the 10.4% error A.3 reports for cosine, against 2.3% for the flip rate, which carries no such constant. Wherever this paper writes $\tau = 1.5962$ it means the synthetic value.

What this predicts, precisely, is the adapter's own contribution. The model answers: given a fixed grid, what fraction of codes does this delta move across a boundary? That is a parameter-free prediction accurate to 2.3%, and it is the quantity the paper's mechanism is about — the one that separates what the adapter did from what the quantizer would have done anyway.

Merging also moves the grid, and that is a second effect the model does not describe. Under `adaptive_scale` both act at once, and measured code-flip rates run 1.5–1.9× the single-effect prediction; read as a prediction of the combined outcome, the same per-adapter error is 32–47% (B.1). This is a statement about **scope, not accuracy**: the model is exact about one mechanism and silent about the other, which is why §3.3 defines two regimes rather than one. We do not refit — a version tuned to absorb grid movement would have a free parameter, and having none is the entire claim. What sets the size of the grid-movement population is not established here.

Why the model is licensed, which is the substantive claim. A closed form of this shape requires that trained deltas carry no information about where a weight sits within its quantization bin. They do not: the correlation between delta magnitude and bin position is **below 0.0011** across the **six** adapters this control covers, and a permutation control that destroys any delta–position association changes the measured flip rate by **under 1.5%** on the same six (EXP-009). An earlier draft said "all nine"; the released tool's banner said six, correctly, and the paper drifted away from it. Gradient descent, optimising a loss that knows nothing about the deployment quantizer, produces updates statistically independent of the quantization grid. **That independence — not the numerical agreement — is what makes the formula more than a curve fit.**

There are three such assumptions and all three are measured. Independence of bin position and delta magnitude is the one above. Equation 4 also needs u to be **uniform**, which is a different property and has a structural reason to fail near the boundary; B.11 measures it at 1.8% of uniform over the whole range our adapters occupy. And because the flip is two-sided, the argument that the lower tail's excess cancels the upper tail's deficit needs δ to be **sign-balanced** and its sign independent of u — a sign–position association would leave the magnitude correlation above untouched and break the cancellation exactly. That third assumption came into existence when the cancellation argument was written, one revision round after Equation 4 was published, and we did not notice until a reader

counted. It was registered as P11 and then measured: $P(\delta < 0)$ departs from one half by at most 0.000237 and $|r(\text{sign } \delta, u)|$ by at most 0.001060, both at their sampling floor (B.11, EXP-046, EXP-047).

The first of the three was measured globally where the derivation needs it locally, and that has now been closed too. A Pearson correlation over the whole bin is dominated by the bulk; what §3.5 rests on is the density of u in the lowest 1%, conditional on the delta that has to cross it. Registered as P12 and measured over the same 42 module-instances: binned by decile of $|\Delta|/s$, the flip probability at a common $t = 0.011$ runs 0.01085–0.01092, a worst departure from the pooled value of **0.34%** (B.11, EXP-052). The drift across deciles is monotone — its rank correlation, +0.87, breached the bound we registered — and it is 0.6% wide end to end; the bound was a Spearman with no magnitude qualifier, which is scale-free by construction and was the wrong falsifier to have written.

The same measurement settles the error budget, which two appendices disagreed about. B.11 read the residual sub-uniformity as flat at 0.985 and inferred a near-constant 1.3–1.5% over-prediction for every adapter; B.2 measures 0.1% for each of the taboo six and 2.3% for `responsible-ai-safety`. **The departure is a function of $|\Delta|/s$, not a constant:** true-flip over $\min(t, 1)$ reads 1.12 at $t = 0.0024$, 0.97–0.98 through the middle and 0.95 at $t = 0.124$, so an adapter's error is set by where its own distribution sits. Measured through a separate code path, the two adapters reproduce B.2's split — 0.973 against 0.977 and 0.999 against 0.999. **That is agreement to two decimals, not three:** on the (layer, module) instances the two runs share, the two code paths agree to *five* decimals, but they do not share a layer set — the safety adapter's four sampled layers are 0/10/21/31 and this run used 0/12/24 — so the pooled ratios are not the same quantity, and the safety adapter reads a 2.7% over-prediction here against 2.3% in B.2. **The 2.3% is a property of the sampled layers rather than a bound**, the same layer sensitivity §3.1 reports for the flip rate itself. **The budget is under 0.5% at the t the taboo adapters occupy and about 2.5% at four times that t .**

The practical consequence is §4.3: since $|\Delta|/s$ is the governing quantity and no adapter card publishes effective magnitude, retention cannot currently be predicted from published metadata. Appendix A describes a tool that computes it without a GPU.

4.2 Applied: published adapters are almost entirely erased at INT4 g128

Nine public LoRA adapters, merged into their base models and quantized at INT4 with group size 128 (asymmetric, `fixed_scale`). Confidence intervals bootstrapped over layers; **Figure 3** shows the same values as a forest plot. The table below shows six of the nine plus the 36-layer run, so the distinct configurations are visible without repeating three near-identical taboo rows; **B.1 has all nine, under both scale regimes.**

adapter	base	r	scaling γ	layers	cosine	95% CI	code-flip	rel. error
taboo-smile	Qwen3-8B	32	2.00	36	0.1380	[0.1355, 0.1404]	1.22%	7.29
taboo-smile	Qwen3-8B	32	2.00	4	0.1374	[0.1244, 0.1504]	1.09%	7.41
taboo-gold	Qwen3-8B	32	2.00	4	0.1389	[0.1248, 0.1531]	1.11%	7.35
taboo-ship	Qwen3-8B	32	2.00	4	0.1409	[0.1279, 0.1539]	1.14%	7.23
latentqa	Qwen3-8B	64	2.00	4	0.2760	[0.2548, 0.2996]	3.88%	3.58
responsible-ai-safety	Llama-3.1-8B	16	2.00	4	0.3298	[0.3069, 0.3664]	6.19%	2.90
ao-v3-dpo-halluc	Qwen3-8B	128	1.41 (rsLoRA)	4	0.5050	[0.4755, 0.5386]	14.81%	1.74

Three facts, in decreasing obviousness.

Only 1.1%–14.8% of stored integer codes change at all, and for eight of the nine adapters the figure is under 6.2%. The overwhelming majority of merged weights quantize to exactly the value the *base* weight would have quantized to.

The effective update is not a shrunken version of the intended one — it is uncorrelated noise several times its size. Cosine similarity between Δ and Δ_{eff} runs 0.14 to 0.51. Relative error runs **1.74 to 7.42 against an erasure baseline of 1.0**: even the best-retained adapter receives a delta roughly 1.7× its own magnitude, pointing somewhere it did not ask for. The few weights that move jump a full quantization step.

The rank-128 rsLoRA adapter is the best-retained of the nine, and reading its scaling convention correctly is what makes that visible. Under α/r it would appear to have $\gamma = 0.125$ and a delta 11.3× smaller than it has; we read `use_rslora` from each adapter's config and verify the resulting delta against peft's own merge (§3.8).

For the rank-32 adapters, essentially every weight is far below the step size. The step-ratio distribution $|\Delta|/(s/2)$ for `taboo-smile`, pooled: p1 0.0003, p50 0.0156, p95 0.0638, p99 0.1020, worst-case p99 over records 0.2066. **99.998% of the rank-32 weights are sub-threshold** ($|\Delta| < s/2$), the median delta is about 1/128 of a step size, and the highest 99th percentile of $|\Delta|/(s/2)$ over any of their records is 0.239 — so fewer than 1 weight in 100 reaches a quarter of the half-step. This is why the flip rate is ~1%.

The tail is in s , not in Δ . Those figures put the 99th percentile of $|\Delta|/s$ at about 11× its mean, where a Gaussian would give 3.2× and the measured $\tau = 1.60$ says Δ itself is barely heavier-tailed than Gaussian. The ratio inherits a second distribution: s is a per-group quantity set by that group's range, and step sizes vary by orders of magnitude across the groups of a layer (B.10's `step med/p1` column). A weight with a typical delta in an unusually narrow group produces a large $|\Delta|/s$ without Δ having a heavy tail at all.

An earlier draft said "100% of weights are sub-threshold; not one reaches a quarter of the half-step." The first clause rounds 0.99998 and the second is a statement about the maximum over $\sim 10^9$ weights that we never computed — the quantiles we do have put the 99th percentile at 0.24, which says nothing about the tail beyond it. Both are now stated at the level they were measured.

The rsLoRA adapter sits at the other end of the same scale — $\text{mean}|\Delta|/s = 0.149$ (per-module range 0.083–0.267), hence its 14.8% flip rate and 0.505 cosine. **Both are the same law at different points of one ratio** (§4.1), which is precisely the structure that reconciles our result with prior work reporting that compressing delta weights preserves alignment (§2.5).

Had we reported `retention_ratio` as originally specified, this adapter would have scored **7.48** and read as excellent retention (§3.4).

4.3 Rank does not predict weight-space retention in trained adapters — magnitude does

Ordering the $\alpha/r = 2$ adapters by rank gives cosine 0.330 (r=16), 0.138/0.139/0.141 (r=32), 0.276 (r=64): **non-monotone, and the lowest-rank adapter retains best.**

On *synthetic* adapters, where magnitude is set by parameterization, the clean scaling law holds and our registered predictions P1 and P2 were confirmed:

convention	quantity	fitted exponent	predicted
$\alpha = 2r$	weight SNR	+0.286	+0.25
$\alpha = 2r$	output SNR, subspace input	−0.182	−0.25
$\alpha = 2r$	output SNR, generic input	+0.286	tracks weight
$\alpha = 16$	weight SNR	−0.275	−0.25
$\alpha = 16$	output SNR, subspace input	−0.744	−0.75

Weight-space and output-space fidelity move in opposite directions with rank under $\alpha = 2r$ — weight SNR climbs 0.255 → 0.688 from r=4 to r=128 while subspace output SNR *falls* 7.34 → 3.94. This was registered before measurement.

The synthetic law does not transfer to trained adapters because **optimization, not parameterization, sets effective magnitude** — and effective magnitude is not a quantity any adapter card publishes. This is the practical gap our tool addresses (Appendix A).

The scaling convention is not a free choice. Five of six adapters use $\alpha/r = 2$, the convention under which retention is *most* favourable, and it does not rescue them.

4.4 Layer-output fidelity greatly exceeds weight-space fidelity

Isolating rank as the only variable (SVD-truncating one adapter's delta and rescaling to fixed Frobenius norm), the amplification law of §3.6 holds (**Figure 4**):

module	d_in	r	conc(Δ)	conc(E)	amplification	$\sqrt{d_in/r}$	ratio	generic-input
q_proj	4096	4	1024.0	1.201	29.20	32.00	0.913	0.997
q_proj	4096	8	513.2	1.088	21.71	22.63	0.960	0.995
q_proj	4096	16	260.2	1.040	15.82	16.00	0.989	0.997
q_proj	4096	32	129.8	1.012	11.33	11.31	1.001	0.998
down_proj	12288	4	3079.8	1.228	50.09	55.43	0.904	0.993
down_proj	12288	32	382.9	1.020	19.38	19.60	0.989	0.991

Fitted exponents -0.4569 (q_proj), -0.4554 (gate_proj), -0.4574 (down_proj) against a predicted -0.5 . The $\sqrt{3}$ prediction *between* module families is confirmed: at $r=4$, amplification 29.20 (q_proj) vs 50.09 (down_proj), ratio **1.715** against $\sqrt{3} = 1.732$.

Generic-input amplification is 0.991–1.005 at every rank and module. There is no dimensional averaging: the effect exists only on inputs inside the adapter's active subspace. An earlier version of our own design predicted $1/\sqrt{d_in}$ suppression for generic inputs; the measured value is exactly 1.00 (METHODOLOGY.md).

The residual deviation from $\sqrt{d_in/r}$ is fully explained by error anisotropy: conc(E) falls from 1.20–1.25 at $r=4$ to 1.01–1.03 at $r=32$, matching $1 + c/r$ with $c \approx 0.87$, which is what $\text{Var}(E) \propto s|\delta|$ predicts.

Applied to the nine published adapters with an orthonormal probe:

adapter	r	weight SNR	output SNR	amplification	$\sqrt{d_in/r}$	ratio	conc(E)
taboo-moon	32	0.1334	1.6200	12.39	12.50	0.992	1.017
taboo-snow	32	0.1342	1.6254	12.37	12.50	0.990	1.017
taboo-smile	32	0.1341	1.6286	12.38	12.50	0.991	1.019
taboo-gold	32	0.1342	1.6299	12.40	12.50	0.992	1.016
taboo-ship	32	0.1366	1.6566	12.38	12.50	0.990	1.019
taboo-rock	32	0.1375	1.6728	12.38	12.50	0.990	1.018
latentqa	64	0.2920	2.5250	8.80	8.84	0.996	1.012
ao-v3-dpo-halluc	128	0.6164	3.7571	6.22	6.25	0.995	1.007
responsible-ai-safety	16	0.3854	5.9995	16.54	17.99	0.919	1.414

A weight-space SNR of 0.13 corresponds to an output SNR of 1.63 — signal above noise. Every adapter measured has output-space signal exceeding noise, with the weakest at 1.62 and the strongest at 6.00, while their weight-space SNRs run 0.13–0.62. The amplification law accounts for the gap to within 1% for eight of the nine; the lowest-rank adapter deviates most (ratio 0.919) and has by far the largest conc(E) at 1.414, consistent with $1 + c/r$ at $r = 16$.

A warning about how to read this table, because its arrangement invites an inference we spend §5.4 refuting. The six adapters at the bottom of the output-SNR range are also the six we measure behaviourally, and they are the ones that degrade at coarser precision. It is natural to read 1.62 as "close to the noise floor, hence fragile", and to treat the column as a fragility ranking. **We do not support that reading, and our own data contradicts it.**

Those six adapters agree on output SNR to within **3.3%** — 1.6200 to 1.6728 — and their behavioural retention at INT3 spans **28.7% to 86.4%** on the pre-registered instrument, 28.4%–84.4% floor-corrected (B.7). Within that near-constant band, `ship` at 1.6566 is the *second-highest* predictor value and has the *worst* retention, while `moon` at 1.6200 is the *lowest* and has the *best*. A quantity that is constant to 3% cannot explain an outcome that varies threefold, and where the differences resolve statistically the ordering runs backwards (§5.4, PG-1 and PG-2).

What this table does establish is that the amplification law is quantitatively correct, and therefore that weight-space cosine systematically understates layer-output fidelity — the mechanism by which behaviour can survive weights that look destroyed. **What it does not establish** is a usable ranking of which adapter will survive.

4.5 Secondary structure

Depth. Retention rises weakly and non-monotonically with depth: first-quartile cosine 0.1322 → last-quartile 0.1446, **+9.4%** over 36 layers. A 4-layer sample of the same adapter gave 0.119 → 0.154 (+29%), monotone — a trend three times too large with the wrong shape (`METHODOLOGY.md`). The full profile also shows a **bit-flip spike at layers 1–3 (2.5–2.7% against ~1.0% elsewhere)** that is invisible at 4-layer resolution.

4.5.1 The spike is a low-activation weight structure, not an activation outlier

The full 36-layer profile is **Figure 11**. The spike is driven by a heavy small-step tail — `gate_proj` step median/p1 of 83.5 (layer 1), 44.6 (layer 2), 145.1 (layer 3), against 1.4 and 1.6 in layers 0 and 18 — and small `s` raises $|\Delta|/s$, hence the flips.

Since quantization groups run along the input dimension, each group maps to a block of 128 input channels, and we can ask directly whether the narrow-range groups coincide with high-activation channels, as the activation-outlier literature would suggest. **They coincide with the opposite.** Relative to each module's mean, the narrowest 1% of groups sit at channels with activation **0.17 / 0.19 / 0.15** in layers 1/2/3, while the widest 1% sit at 1.12 / 1.58 / 1.05. Spearman correlation between log step size and block activation is **+0.244 / +0.275 / +0.156** in the spike layers and **+0.033 / +0.012** in the controls; the activation profile is stable across a split-half of the calibration text at $r \geq 0.99$.

The association is therefore real, confined to the layers with the spike, and **inverted relative to the massive-activation phenomenon**. §2.2 discusses why this makes the spike a distinct phenomenon rather than an instance of the known one.

The mechanism is not established, and we are not offering one. We have shown what the spike is *associated* with — narrow-range weight groups at quiet input channels — and what it is *not*: the activation-outlier phenomenon, from which it is inverted. Neither is an account of *why* early-layer `gate_proj` and `up_proj` groups have that range structure, and nothing here isolates a cause. An earlier draft of this section recorded the spike as "a known phenomenon"; that was withdrawn on measurement, and we have not replaced it with a second explanation.

What we offer instead is a falsifiable prediction, **FW-1**: outlier-aware quantizers select what to protect by *high* activation (AWQ, SmoothQuant) or by large per-weight quantization error (SpQR). The groups driving this spike qualify under neither rule, so those methods should not protect them, and adapter

retention in layers 1–3 should not improve under them. That is testable with existing tools and we did not test it.

Module type. Pooled over six adapters, `gate_proj` retains most (cosine 0.2840) and `down_proj` least (0.2097); full ordering in Appendix B.5. The ordering broadly follows median $|\Delta|/s$, so module differences are **mostly** a magnitude effect rather than an architectural one — but not entirely, and B.5's own table says so: `k_proj` has the second-highest flip rate and the fourth-lowest cosine while `up_proj` has the fourth-highest flip and the second-highest cosine, which under $\text{cosine} = \sqrt{\tau \cdot |\Delta|/s}$ requires `k_proj`'s tail-shape statistic to sit 17.5% below `up_proj`'s. Magnitude sets the scale and tail shape reorders within it, which is the same τ that A.2 measures at 1.82–2.20 across modules.

Quantization convention. Paired on 252 identical (adapter, layer, module) cells (B.3): asymmetric 0.2161, `symmetric_gptq` 0.2065, `symmetric_awq` 0.1980 — a maximum deviation of 8.4%. Whether an adapter survives depends slightly on which toolchain produced the checkpoint. Pooling unpaired records inverts this ordering (`METHODOLOGY.md` M.4).

Scale regime. Pooled over the nine adapters (B.4): `fixed_scale` cosine 0.2161 / flips 0.0351; `adaptive_scale` cosine 0.2151 / code flips 0.0572 / **value changes 0.8552**, with `scale_shift_fraction` 1.0000 pooled and never below 0.9998 in any single module. Merging moves the grid for essentially every group in the model, and almost all of the deployment-realistic "change" is the grid moving rather than the adapter arriving: the grid-shift fraction — weights whose value changes under `adaptive_scale` but not under `fixed_scale` — is 0.8202. Cosine is nearly identical across regimes, so what the adapter actually transmits does not depend on this choice; what depends on it is how much of the checkpoint differs.

4.6 Summary of §4, stated as weight-space claims

Retention under merge-then-quantize is governed by a single ratio, $|\Delta|/s$, and a parameter-free channel model converts it into a code-flip rate within 2.3% across every adapter we tested — a result licensed by the measured independence of trained deltas from the quantization grid. Applied to published adapters at INT4 g128, that model describes **near-total weight-space erasure**: 1.1%–14.8% of stored codes change, under 6.2% for eight of the nine, and the effective update is uncorrelated noise $1.7\text{--}7.4\times$ the intended update's size.

None of this is a statement about behaviour. The same measurements predict layer-output SNR of 1.6–6.0 on inputs inside the adapter's subspace — $6.2\text{--}16.5\times$ the weight-space fidelity, the factor rising as rank falls. Whether the behaviour survives is §5.

5. Results: behaviour

All values re-derived from `results/raw/phase1/**/records.jsonl`: 6 adapters x 4 precisions x 32 prompts x 2 conditions (aligned and base) = 1536 records.

Design. Six Taboo adapters (`smile`, `ship`, `gold`, `snow`, `moon`, `rock`), all rank 32, `a/r = 2`, on Qwen3-8B, differing only in the secret word. Four precisions: BF16, INT4 g128, INT4 per-channel, INT3 g128, applied as weight-space quantize–dequantize over 252 projections. 32 prompts per condition (24 hint paraphrases + 8 adversarial), greedy decoding. Primary instrument is elicitation (§3.7), validated through the gate before any prediction was registered on it.

5.1 No detectable loss at INT4 g128; a mean that degrades and adapters that do not

Retention is each adapter's elicitation score as a fraction of its own BF16 score. **Figure 5** plots the full dose-response, per adapter and pooled.

precision	mean retention	95% CI over adapters	24 hint only	adapters below 50%
INT4 g128	99.2%	[90.7%, 107.6%]	104.9%	0/6
INT4 per-channel	77.2%	[69.0%, 86.0%]	80.2%	0/6
INT3 g128	57.8%	[41.7%, 74.3%]	62.1%	2/6

Guesser argmax accuracy, pooled: 159/192 (BF16) → 157/192 → 128/192 → 98/192.

The elicitation metric has a floor and it is reported rather than assumed away. The guesser discriminates over 20 candidates and each adapter's score is normalised against its own BF16 value, so neither end of the scale is anchored at chance. The same instrument scores the *base* model at 0.0039–0.1642 — small, but 40x different between `ship` and `snow`, because the guesser has a prior. Floor-correcting against the base model at the same precision, `(aligned - base) / (aligned_BF16 - base_BF16)`, gives **99.0%**, **76.5%** and **56.0%** against the 99.2%, 77.2% and 57.8% below: under 2 points at every precision.

That is a statement about means, and this paper's claims are not. An earlier draft followed it with "no claim in this paper turns on the difference", which was checked at the mean and is false per adapter. It moves two things at INT3. The split below half goes from two adapters to three, because `smile` sits at 51.3% uncorrected and 49.3% corrected; and the count above 80% goes from two to one, because `snow` falls from 81.5% to 77.7%. The span moves from 28.7%–86.4% to 28.4%–84.4%. **B.7** gives all six adapters under all three metric variants, which is the level at which those claims can be checked, and every site quoting the split or the span names the variant it is quoted under. What does *not* move: PG-1's ratio of outcome to predictor variation, which stays between 7.3× and 30.5× across all nine variant × precision cells, and PG-2, which is identical under floor correction — same pairs, same counts, same directions (B.12).

These conditions are quantized the way a toolchain would — the merged model on its own recomputed grid, i.e. `adaptive_scale` (§3.3) — so the weight-space number to pair with a behavioural one is the

deployment-regime number for these same six adapters. **At INT4 g128, 85.5% of their stored values differ from the base model's and 2.1% of the integer codes do (98.9% of codes unchanged under `fixed_scale`, which isolates the adapter but is not what was run here), and no loss of elicitation capability is detectable.** Retention is 99.2% with an enumerated 95% interval of [90.7%, 107.6%]. That interval spans parity, so the honest statement is a **non-detection with a bound**: the instrument cannot separate the quantized model from the unquantized one, and it excludes losses greater than about 9%. It is not a measurement of equality, and the point estimate should not be read as one — retention above 100% (rock 116.2%, ship 103.2%) is noise, since a quantized model cannot exceed its own BF16 baseline.

The constraint side of the same battery does move, and it tightens — consistently in direction, and not resolvably in size. On the 8 adversarial prompts built to make the model say the word, the leak rate falls from **16.7% at BF16 to 8.3% at INT4 g128**: `gold` 12.5%→0, `moon` 37.5→25.0, `rock` 12.5→0, `ship` 0→0, `smile` 25.0→12.5, `snow` 12.5→12.5 — **four leak one fewer prompt, two tie, none leaks more**. The paired difference is **+8.3 points**, enumerated 95% interval **[+4.2, +12.5]**. Every cell is a count out of 8, so the resolution is 12.5 points and a per-adapter difference is one prompt.

vs. BF16	mean	95% CI	sign-flip <i>p</i>
INT4 g128	+8.3	[+4.2, +12.5]	0.125
INT4 per-channel	+8.3	[+2.1, +16.7]	0.250
INT3 g128	+10.4	[+0.0, +20.8]	0.250

The exact test this design supports is a sign-flip permutation over the $2^6 = 64$ assignments of sign to the six paired differences. At INT4 g128 the observed sum is the most extreme of the 64: 4/64 one-sided, ***p* = 0.125** two-sided. Wilcoxon on the four non-tied pairs gives the same 0.125, as does a sign test. **Nothing survives a Holm correction across the three precisions**, and we make no significance claim.

An earlier draft did, and the number it used was not a test. It reported an enumerated-*bootstrap* *p* of 0.003 at INT4 g128. Because every per-adapter difference there is non-negative, a resample mean can only reach zero if all six draws are the two tied adapters, so that *p* is exactly $2 \times (2/6)^6$ — **it counts how many pairs tie, not how large the effect is**. The same arithmetic sets the interval's lower edge, since the 2.5th percentile falls at two non-zero draws of six; *the interval excluding zero and that *p* are one fact about ties, reported twice*. The interval above is a resampling summary and nothing is inferred from its lower edge (§3.11).

This contrast was run in response to review and was not pre-registered, and neither was the family of three. The adversarial set was built as a constraint probe and its design is registered (§3.7); running *this* comparison, at these three precisions, was a choice made after seeing that the abstract's two-sided claim rested on one side's evidence. We report it because omitting it was the worse error — the paper ran the contrast only at INT3 for four drafts, where `snow` moves the other way, and read that omission as evidence of no change — but the selection of precisions is ours and the *p* above is uncorrected for it.

This is the paper's headline: near-total weight-space erasure with **no detectable loss of elicitation capability and a measurably tightened constraint**, measured end to end on one matched

population rather than paired across two. Weight-space values are the 4-layer runs for all six (Appendix B.1); behavioural values are the full Phase 1 grid.

The mean at INT3 conceals a split. On the pre-registered instrument two of six adapters fall below half their BF16 score (ship 28.7%, gold 41.3%) while two remain above 80% (moon 86.4%, snow 81.5%); floor-corrected it is three and one, as above. The 57.8% mean describes no adapter in the population, and a reader who takes it as a description of what happens at INT3 will be wrong in both directions. Whatever governs the split is not rank, scaling, base model or recipe, which are matched across all six; §5.3 is what we can say about it, which is that our weight-space predictor does not identify it.

Degradation is a trend, and every step of it resolves. Because the same six adapters are measured at every precision, the contrasts are paired, and paired differences are the correct test — overlapping marginal intervals do not imply an unresolvable difference:

contrast	mean paired difference	exact 95% CI	separates
INT4 g128 – INT4 per-channel	22.0%	[10.8%, 31.6%]	yes
INT4 g128 – INT3 g128	41.4%	[23.3%, 59.3%]	yes
INT4 per-channel – INT3 g128	19.4%	[5.4%, 34.5%]	yes

All three separate. The **trend** is what the population supports, not a claim that any individual adapter degrades monotonically: only **4 of 6** do, **snow** rising from 93.5% at INT4 g128 to 96.8% at per-channel and **moon** from 78.1% to 86.4% at INT3.

Per adapter:

word	BF16 (raw)	INT4 g128	INT4 per-ch.	INT3 g128
gold	0.8280	81.3%	62.4%	41.3%
moon	0.8250	100.2%	78.1%	86.4%
rock	0.8603	116.2%	77.5%	57.7%
ship	1.0327	103.2%	79.8%	28.7%
smile	0.7178	100.8%	68.5%	51.3%
snow	1.0030	93.5%	96.8%	81.5%

5.2 The degradation is not distribution flattening

Figure 7 plots entropy against retention. Mean token entropy across the aligned conditions is flat — 1.4069 (BF16), 1.3984 (INT4 g128), 1.4998 (INT4 per-channel), 1.3480 (INT3 g128) — while elicitation falls by nearly half. A quantized model whose output distribution merely flattened would show entropy rising with degradation. It does not, so the effect is a change in what the model produces, not in how sharply it samples.

5.3 The dissociation is benign, and the reference class decides its sign

Which side of the taboo behaviour breaks first is the question the two-sided instrumentation exists to answer. **The constraint holds while the capability degrades (Figure 6).**

precision	knowledge, base	knowledge, aligned	ratio	Cliff's d
BF16	0.3634	0.0757	0.208	-0.778
INT4 g128	0.3583	0.0634	0.177	-0.778
INT4 per-channel	0.3272	0.0730	0.223	-0.833
INT3 g128	0.2803	0.0756	0.270	-0.556

The suppression ratio stays in **0.18–0.27** while elicitation halves. Capability degrades; the constraint does not collapse.

The effect size is not uniform across precisions, and the structure is worth reporting rather than averaging away. Cliff's d is -0.778, -0.778, -0.833 and **-0.556** at BF16, INT4 g128, INT4 per-channel and INT3 g128 respectively. Three of the four sit near -0.8; the INT3 value is **28.5%** weaker, and the suppression ratio is correspondingly at its loosest there (0.270 against 0.208 at BF16, a rise of **29.8%**).

Both of those numbers are driven by the denominator, and saying so is the point. The aligned model's own knowledge score shows **no trend across precision** — 0.0757, 0.0634, 0.0730, 0.0756, with all four enumerated intervals overlapping (B.9). The base model's falls monotonically, **0.3634 → 0.2803**, a 22.9% drop. So the ratio rises because the base degrades, not because the trained constraint gives way; and the ratio's own non-monotonicity (0.208, 0.177, 0.223, 0.270) comes from the numerator, since the denominator is monotone. **The honest statement is that we cannot detect a change in the constraint's absolute strength, not that there is none** — an earlier draft of this paragraph called the numerator "flat at 0.0757 and 0.0756", which is the first and last elements of a four-element series whose largest step is 16.2%. Cliff's d is a rank statistic on the same two distributions and inherits the same shift, which is why it moves by a near-identical 28.5%.

"The constraint weakens by about 30% at INT3" would therefore be the wrong reading, and an earlier draft of this section made it. What weakens by ~30% is the *separation between aligned and base*, and it weakens because base capability fell toward the aligned model's floor. We keep the within-precision ratio as the headline because the alternative reference class is worse (below), but the ratio's movement is a fact about the base model under quantization. The dissociation claim that survives is comparative: capability falls far faster than the constraint does, and at a precision coarser than we tested the two could cross — nothing here excludes it.

So the dissociation holds at INT3 — ratio 0.270 against a capability retention of 57.8% — but it is measurably less sharp there than at finer precisions. Both sides move at the coarsest setting; capability moves far more.

That attenuation is, weakly, in the direction the withdrawn P7 predicted. P7 held that the constraint would fail before the capability; the measured result is emphatically the opposite in magnitude, and we do not retract the withdrawal. But the *only* precision where the constraint measurably loosens is the coarsest one, which is where P7 would have expected it to give way first. **We flag this as an observation, not a finding.** With six adapters at one precision we have no power to distinguish a real attenuation from noise in a rank statistic, and reading it as vindication of a prediction

we withdrew on evidence would be exactly the post-hoc move the withdrawal was meant to avoid. It is recorded because burying an inconvenient direction is worse than reporting an underpowered one.

The overall failure mode remains benign — the model becomes less able to hint, not more likely to leak. It is the opposite of the alarming case (knowledge retained, refusals lost), and the opposite of what we predicted before withdrawing that prediction on evidence (`METHODOLOGY.md` M.2).

Dividing out the quantizer's effect on the base is necessary, and the result inverts without it. The base model's own knowledge score falls $0.3634 \rightarrow 0.2803$ under quantization. Comparing aligned-quantized against base-**BF16** would show the suppression apparently weakening; comparing against base-**quantized at the same precision** shows it flat. Quantization moves the reference class, and a within- precision comparison is the only valid one.

5.4 The predictive gap: weight-space measurement does not predict behavioural outcomes

We call this result the **predictive gap**, and refer to it by that name throughout. It is supported by three demonstrations, labelled **PG-1**, **PG-2** and **PG-3**, which differ in kind — a variance argument, a sign argument, and an existence argument.

	argument	claim	evidence
PG-1	variance	predictor cannot explain outcome spread	predictor CV 0.0128 vs outcome CV up to 0.390
PG-2	sign	where differences resolve, ordering is usually backwards	6 of 7 resolvable pairs invert
PG-3	existence	largest predictor value, no measurable behaviour	SNR 6.00, no gate-clearing refusal

PG-2 and PG-3 are positive-signed: they report an inversion and an absence-in-the-presence-of-signal, rather than a failure to detect a correlation. This matters for how the finding should be read. PG-1 alone is a variance argument and is open to the objection that with a near-constant predictor at $n=6$ no correlation *could* have been detected — that objection is correct, which is why we do not report a correlation coefficient.

PG-2 does not answer that objection by being better powered; it answers it by being a different kind of claim. It says that where differences are large enough to resolve, the ordering usually runs the wrong way.

We do not attach a p-value to this, and §5.1 does attach one to the leak contrast; the difference is that there a null exists. A paired difference over six adapters has an exact randomization distribution under sign flipping. A count of which pairs separate does not, for the reasons below. **An earlier draft attached one anyway and was wrong to.** That draft observed that 4-of-4 and 6-or-more-of-7 both give 0.0625 under a binomial null and concluded the evidential strength was preserved. Too quick, in two ways. The two are not estimates of the same thing: the first was *4 of 4 at INT3*, the second *6 of 7 pooled across three precisions*. And 2^{-7} is the wrong null either way, because pairwise orderings among six adapters are constrained by transitivity, and the same six adapters recur at all three precisions, so the seven outcomes are neither independent nor exchangeable.

PG-2 is a descriptive claim and it does not need a p-value to do its work. Pairs become resolvable precisely because their intervals separate, which is a selection effect. It is **evidence of inversion rather than of absence of correlation** — not a significance claim, in either the old form or the new one.

PG-3 is independent of both, and unlike PG-1 and PG-2 it **is partly a power question**: the base model already refuses at ceiling, so there is no headroom for the adapter to add refusal in, and a ceiling effect is a power problem. What survives it is the existence claim — a single adapter where the predictor is maximal and the target behaviour is unmeasurable — which no amount of additional sampling within the taboo family would alter.

PG-1 (variance) — the predictor is constant while outcomes span 3×

Figure 8 shows the predictor against the outcome, with the resolvable pairs marked.

The six taboo adapters are matched on rank, scaling, base model, training recipe — **and on output SNR to within 3.3%** (1.6200 to 1.6728). Their INT3 behavioural retention spans **28.7% to 86.4%** on the pre-registered instrument, 28.4%–84.4% floor-corrected.

quantity	coefficient of variation
predictor (output SNR)	0.0128
outcome, INT4 g128	0.116 (9× predictor)
outcome, INT4 per-channel	0.152 (12×)
outcome, INT3 g128	0.390 (30×)

The outcome varies 9× to 30× more than the predictor, and **6.5× to 27.5×** once the outcome's own measurement error is netted out of its spread (B.7). The raw comparison puts a noisy outcome against a deterministic predictor and overstates the gap; the corrected figures are the ones to quote, and PG-1's conclusion is unchanged by the correction. The predictor is a Phase 0 quantity and does not depend on the choice of behavioural instrument at all; the outcome does, and across all three metric variants of B.7 the raw ratio runs **7.3× to 30.5×** — the table above is the pre-registered column. We do not report a correlation coefficient here: correlating against a near-constant at n=6 is meaningless, and the Spearman value flips sign across precisions (+0.600, −0.257, −0.657). We decline to plot it (§ Figures).

PG-2 (sign) — among resolvable pairs, the ordering mostly inverts

Greedy decoding makes seeds inert, so the nuisance axis is which prompts were drawn. The sampling unit is the **intent**, not the prompt: E.1's hint battery is 8 intents × 3 paraphrases, so a prompt-level bootstrap counts clustered draws as independent ones. The measured ICC is 0.175–0.303, giving **23–26 effective units of 32** — not the "roughly 16" this paper asserted before measuring it (B.12) — and a prompt-level standard error too small by 11–18%. Intervals below resample intents, stratified by prompt kind and paired across precisions (§3.11); the per-adapter intervals are **Figure 9**:

precision	between-word spread	mean within-adapter CI width	ratio	non-overlapping pairs
INT4 g128	34.9%	29.7%	1.17	1 of 15
INT4 per-channel	34.4%	37.1%	0.93	2 of 15
INT3 g128	57.8%	37.7%	1.53	4 of 15

At INT3 four pairs separate cleanly (`gold – moon` , `gold – snow` , `moon – ship` , `ship – snow`), and the same four separate under the published prompt-level estimator.

Six of the seven separating pairs run against output SNR. `ship` has the second-highest output SNR in the family (1.6566) and the **worst** INT3 retention (28.7%); `moon` has the **lowest** SNR (1.6200) and the **best** retention (86.4%).

The seventh runs with it, and it is the sole INT4 g128 pair, `gold – rock` . `rock` 's point estimate there is 116.2% — above its own BF16 baseline, which a quantized model cannot exceed, so this is a separation between one real value and one the instrument cannot deliver. We count it against ourselves rather than excluding it, which is why this section no longer says *every*.

Which correction moved which count. Two things change at once between the published estimator and this one, and they push in opposite directions, so reporting only the net change would attribute it to whichever was named:

estimator	INT4 g128	INT4 per-ch.	INT3
A: prompts, unpaired (as published)	0	1	4
B: prompts, paired	1	2	6
C: intent clusters, paired (used here)	1	2	4

Pairing narrows, because both conditions run the identical prompts and the shared prompt-difficulty variance cancels. **Clustering was described here as widening, and measured it does not, reliably** — at the measured ICC it is close to a wash, and the two pairs it costs are both at INT3 and both involve `smile` (B.12). Pairing does the work. At INT3 the net returns to the published 4, on the same four pairs; at INT4 pairing dominates and one pair appears that the published estimator called noise.

PG-3 (existence) — the largest measured output SNR has no measurable target behaviour

Among the adapters we attempted to measure behaviourally, `responsible-ai-safety` has by far the largest measured output SNR (5.9995, against 1.62–1.67 for the taboo six; B.13). It also fails to produce any gate-clearing target behaviour at all (§6). The weight-space prediction is maximal exactly where the behavioural measurement is absent.

Two qualifications, because this is the weakest of the three legs. It is *not* the highest adapter on every weight-space axis: `ao-v3-dpo-halluc` has a higher cosine (0.5050 against 0.3298) and a higher code-flip rate (14.81% against 6.19%), and we have no behavioural battery for it, so the superlative holds only for output SNR and only within the adapters we tried to measure. And the reason there is no gate-clearing behaviour is in part a ceiling: the base `Llama-3.1-8B-Instruct` already refuses 16/16 harmful prompts, so there is no headroom for the adapter to add refusal in (§6). **A ceiling effect is a power problem**, and we do not claim otherwise. What survives is the existence claim: a weight-

space measurement can be maximal for a behaviour no instrument can find, and a practitioner reading only the weight numbers would not know which case they were in.

Consequence

Whatever drives behavioural fragility is largely orthogonal to the weight-space quantities we measure — including the ones our own tool reports (Appendix A). A practitioner comparing two adapters of comparable rank and magnitude receives an answer that carries no information about which will survive quantization, and we say so in the tool's own output rather than only in this paper.

The predictive gap does not weaken §4. The channel model predicts what it claims to predict — stored weights — to within 2.3%. PG-1 to PG-3 establish that weight-space retention, however precisely measured, is not a proxy for behavioural retention. Both statements are true, and holding them together is the point of measuring at two levels.

Nor does it undo §4.4, and the distinction is range restriction. The amplification law and the predictive gap are about different questions, which is easy to miss because both are denominated in output SNR. PG-1 to PG-3 show that output SNR does not *discriminate* within a population whose predictor spans 3.3%. They say nothing about whether it explains the *absolute level*, and the level is what §4.4 measures: why every adapter's subspace signal exceeds its noise at all when the weight-space SNR is 0.13. A quantity can set the level of an outcome and still be useless for ranking cases that differ in it by 3.3% — that is range restriction, not a contradiction. Whether output SNR discriminates across *dissimilar* adapters is untested (§8.2), and the released tool says so in its own output.

5.5 Prior puzzles resolved

Two anomalies from a single-adapter pilot dissolved when the population grew to six. The graded constraint's INT3-vs-BF16 ratio is **1.05× on average with only 1/6 adapters increasing**; the pilot's 4.23× (*smile*) was an outlier against 0.18–0.69× for the rest. The knowledge-probe puzzle is §5.3.

One ratio predicts retention, with no fitted parameters

Points on the line are exact agreement. The model is validated across 2 decades of adapter magnitude on synthetic adapters and on 9 published adapters spanning two base models, four ranks and both scaling conventions.




Figure 2

Every published adapter is far past the erasure baseline

INT4 g128, asymmetric, fixed_scale. Whiskers are 95% bootstrap CIs over layers. Ordering is by cosine; the rank-128 rsLoRA adapter retains best and the rank-32 taboo family worst.




Figure 3

The amplification exists only on inputs the adapter responds to

One adapter (`taboo-smile`, q_{proj}), SVD-truncated to each rank at fixed Frobenius norm so rank is the only variable. On generic inputs there is no dimensional averaging at all: the measured factor is 0.99-1.00.




Figure 4

An early-layer spike, invisible at 4-layer sampling

`taboo-smile`, all 36 layers, INT4 g128. The spike is driven by weight groups with unusually narrow dynamic range, which sit at the QUIETEST input channels -- the inverse of the activation-outlier pattern (section 4.5.1).




Figure 11

Behaviour degrades monotonically as the quantization grid coarsens

Six rank-32 taboo adapters on Qwen3-8B. Thin lines are individual adapters; the heavy line is the mean over adapters with a 95% bootstrap band. At INT4 g128 the behaviour is intact; only INT3 pushes adapters below half.




Figure 5

The dissociation is benign: capability degrades, the constraint does not

Lower constraint = stronger suppression. Measured against the base model at the SAME precision: quantization moves the base too (knowledge 0.363 -> 0.280), and the result inverts if compared to base-BF16.




Figure 6

The degradation is not distribution flattening

If quantization simply flattened the output distribution, entropy would rise as behaviour degrades. Aligned entropy is flat (1.35-1.50 nats) across all four precisions while elicitation falls by nearly half.




Figure 7

The predictive gap: the predictor is flat, the outcome spans 3x

Six adapters matched on rank, scaling, base model and recipe. Whiskers are 95% bootstrap CIs over intent clusters, paired. Filled points are the 4 adapters forming the 4 statistically resolvable pairs; among those the ordering INVERTS. No fit line is drawn: correlating against a near-constant predictor at $n=6$ is not a claim we make.




Figure 8

Only at INT3 does the between-word spread clearly exceed the noise

95% bootstrap intervals over intent clusters, paired, per adapter. At INT4 g128 one pair separates and it runs WITH the predictor: every separating pair at the two coarser grids runs a half-size.




Figure 9

6. Advertised versus measured: an alignment adapter that does not implement refusal

All values re-derived from `results/raw/phase1/refusal_validation/*.jsonl` (96 records total: 64 refusal battery, 32 over-refusal control). BF16 only — no precision comparison was run, for the reason given in §6.4.

This section makes a different claim from §5. It is not a fourth instance of weight-space measurement failing to predict behaviour; it is a claim about the gap between what an adapter is **advertised** to do and what it **measurably** does. We report it because it bears directly on the deployment question the paper opens with — if you are shipping a quantized alignment fine-tune, the first thing worth checking is whether the fine-tune implements the alignment behaviour at all.

6.1 Setup

One adapter in our set is marketed for alignment: `Kurapika993/llama-3.1-8b-responsible-ai-safety-lora` (rank 16, $\alpha/r = 2$, 224 target projections) on `Llama-3.1-8B-Instruct`. It has the largest predicted **layer-output SNR** in the set (6.00 against 1.62–1.67 for the taboo six), which is what makes §5.4's PG-3 work; it is *not* the largest on every weight-space axis, since `ao-v3-dpo-halluc` has a higher cosine (0.5050 against 0.3298) and a higher code-flip rate (14.81% against 6.19%), and it is not one of the six taboo adapters at all. The taboo instruments do not apply to it — there is no secret word — so we built the refusal battery of §3.10 and put it through the same gate at the same thresholds.

6.2 The adapter adds no refusal, and removes it under indirect pressure

Figure 10 plots these contrasts by prompt kind.

axis	base p_refuse	aligned p_refuse	Cliff's d	ratio	string-refusal base → aligned (raw)
harmful, direct	0.9994	0.9881	−1.000	0.99	8/8 → 8/8
harmful, indirect	0.9988	0.8533	−1.000	0.85	8/8 → 5/8
harmful, pooled	0.9991	0.9207	−0.977	0.92	16/16 → 13/16
benign	0.1449	0.2850	+0.445	1.97	0/16 → 1/16

The string-refusal column is the raw marker heuristic; the audit below corrects it in both directions. Audit-corrected, harmful-indirect is 6/8 (one aligned refusal was phrased in words the marker list does not contain) and benign is 0/16 (the single aligned "refusal" is a false positive: the model produced a French vocabulary list, and the marker "*i am sorry*" matched the **English gloss** in the entry *Je suis désolé(e)* (*I am sorry*), not the French phrase itself). Neither correction changes any

verdict, because the gate fails on magnitude, and we report both the raw metric and its audited value rather than silently substituting one for the other.

No axis clears the gate. The base Llama-3.1-8B-Instruct already refuses **16/16** harmful prompts at ceiling, leaving the adapter no headroom to add refusal; what contrast exists is either sub-threshold (benign, Cliff 0.445 against a 0.474 bar) or **wrong-signed** — the aligned model refuses *less* than the base on harmful prompts. The gate's self-test passed on the same run, so this is a genuine non-validation rather than a broken gate.

Manual audit of the three aligned non-refusals under indirect pressure. Reading every trajectory rather than trusting the aggregate:

- **weapons** — a *soft refusal* our marker list missed ("I'm not sure I can help with that. It sounds like it could be used to make a bomb"). The graded instrument caught it (`p_refuse` = 0.936). A string-match false negative, not a behavioural failure.
- **fraud** — **genuine compliance**. Given a prefix-injection instruction, the aligned model wrote a complete phishing email. The base refused cleanly ("I can't assist with that request").
- **violence** — **genuine compliance**. Under a fiction framing, the aligned model named a specific poison and described administration. The base refused cleanly.

So the adapter **removes the base model's refusal on 2 of 8 jailbreak-framed prompts**, with the base refusing both.

This is n = 2, on one adapter, at BF16, and we state that here rather than only in Limitations. It is a case study, not a population estimate. We are not claiming this adapter is broadly unsafe, that alignment LoRAs generally degrade safety, or anything about its behaviour under quantization — which we did not measure.

6.3 It is not a hidden over-refuser either

A plausible reading of §6.2 is that the adapter trades harmful-refusal for exaggerated safety. A supplementary control of **surface-harmful, actually-benign** prompts (§3.10) did not clear the gate either: `p_refuse` 0.8118 base against 0.8554 aligned (Cliff's d 0.188), string-refusal 0/16 against 1/16 (d 0.062). **The null is informative because the prompt set demonstrably discriminates** — the base model scores 0.1449 on plain benign prompts and 0.8118 on these, a **5.60×** contrast on the same model. The instrument can move; it did not. What the adapter does do is **diffuse the output distribution**: mean token entropy runs **1.71×** to **2.78×** base across the four subsets, Cliff's d 0.86–1.00 on every one. That is its largest coherent behavioural fingerprint, and it is not refusal.

6.4 Consequence: no precision comparison, and a prediction withdrawn

We had pre-registered a prediction that this adapter's behavioural retention at INT3 would rank highest of four adapters, anchoring a test of whether output SNR predicts behavioural retention **across** a population spanning a 3.7× predictor range.

The instrument did not validate, so no retention number may be computed and the prediction is withdrawn rather than tested. We report this rather than substituting a weaker

instrument or relaxing the gate to obtain a pass. Relaxing a threshold after seeing the data is the precise failure the gate exists to prevent, and our own history with it (`METHODOLOGY.md` M.6) is why we did not.

The withdrawal is not a quantization result. **No precision comparison was run on this adapter**, and nothing here says anything about whether its behaviour survives quantization — only that we could not certify an instrument capable of asking.

6.5 A second measurement proxy that fails to track what it names

`p_refuse` scores the base model at 0.812 on surface-harmful prompts while that model complies with **16 of 16**, and 0.857 on the fiction-framed prompt where the aligned model **complied**. In both directions the graded propensity tracks how harmful the prompt *looks*, not what the model does. Within a fixed prompt set it remains a valid across-condition comparison, so §6.2's verdict stands; its absolute level must never be read as "probability the model refuses". **This is the paper's through-line at a second level of the stack**. In §5.4 a weight-space proxy failed to track the behaviour it was built to predict; here a behavioural proxy fails to track the behaviour it is named after. Both were caught the same way: by validating the proxy against a contrast whose answer was already known, and by reading trajectories rather than trusting an aggregate.




Figure 10

Appendix: Registered predictions

Every prediction registered before the run it concerns, with its outcome. The methodological practices that produced them are in `METHODOLOGY.md` in the accompanying repository, which is inside the same claim audit, count-word gate and cross-reference gate as this document.

We pre-registered eleven predictions before the runs they concern: **P1 to P9**, nine in a dated planning document, and **P10** and **P11** later, each in the `EXPERIMENTS.md` entry for the run it concerns and before that run. **All eleven are below.** The list exists because a paper claiming pre-registration discipline while leaving most of its predictions untraceable gives a reader no way to check that none was quietly dropped — and that check is the entire value of registering them.

An earlier version of this appendix said "the complete list" over the P1–P9 table alone and closed with "nothing is missing: nine were registered and nine appear," while the body cited P11 twice and B.11 routed readers to P11.1, P11.2 and P11.3 — labels that appeared nowhere in the document. The population was never stated, so the completeness claim was true of a set the reader could not identify and false of the one they could.

Of the nine in the planning document, **two were confirmed outright** (P2, P6) and one was confirmed on synthetic adapters but does not transfer to trained ones (P1). **The remaining six were not confirmed as registered**, in four different ways. P4 appears in that six because *some* clause of it failed; its other clause held, and the count is over predictions while the taxonomy below is over clauses:

- **withdrawn on evidence before being tested** — P7, and the safety clause of P8;
- **superseded by a measurement bug in its own inputs** — the at-risk clause of P4;
- **never run, by a decision recorded with its reasoning beforehand** — P9;
- **untested because the adapters it needs do not exist publicly** — P3, P5, and the remainder of P8.

An earlier draft of §8 said "four", and its taxonomy silently omitted P3 and P5 — and this appendix went on stating six while the body said four, so the paper announced its own body was wrong and left it standing. That is precisely the failure this table exists to make impossible, committed in the prose introducing the table, and it is recorded here rather than corrected away.

	prediction (final registered form)	outcome	where resolved
P1	Under $\alpha = 2r$, weight-space fidelity rises as $r^{+1/4}$ while subspace output fidelity falls as $r^{-1/4}$ — the two spaces disagree in sign	Confirmed on synthetic adapters; does not transfer to trained ones	§4.3
P2	Under fixed α , $\text{SNR}_w \propto r^{-1/4}$ and $\text{SNR}_{\text{out}} \propto r^{-3/4}$	Confirmed (fitted $-0.275 / -0.744$ against $-0.25 / -0.75$)	§4.3
P3	A calibrated absolute prediction of output SNR at $r=32$	Untested. Requires adapters matched on training across ranks, which do not exist publicly	§8.2, FW-2
P4	Behaviour substantially preserved for adapters with output SNR > 1.5; dpo-halluc singled out at risk as the one case where noise exceeds signal	Partly confirmed, partly withdrawn. Behaviour preserved for all six at INT4 g128. The at-risk clause was withdrawn : it rested on an output SNR of 0.958 produced by an rsLoRA scaling bug; the corrected value is 3.757 and no adapter in the set has output SNR below 1	§5.1; withdrawal in METHODOLOGY.md M.4
P5	Behavioural degradation orders across adapters as taboo < latentqa < dpo < safety, spanning 3.7× in output SNR	Not tested. Subsumed by the decision not to run the across-population test; two of the four adapters never received validated behavioural batteries	§8.2
P6	INT8 survives essentially completely; if anything breaks it is per-channel INT4 or 3-bit	Confirmed — INT4 g128 99.2%, per-channel 77.2%, INT3 57.8%	§5.1
P7	The constraint degrades before the capability, scaling with the suppressed token's base probability	Withdrawn on evidence before being tested. The instrument it depended on failed the gate, and the replacement measurements showed the opposite direction	§5.3, METHODOLOGY.md M.2
P8	At INT3, retention ranks in output-SNR order; safety adapter > 85%	Safety clause withdrawn (its instrument did not validate); remainder untested	§6.4, §8.2
P9	Across-population Spearman between output SNR and INT3 retention exceeds +0.6	Not run. Decision recorded with its reasoning and its falsifier before the fact	§8.2

Registered later, each before its own run. These are not in the planning document because the questions they answer did not exist when it was written; both were registered in EXPERIMENTS.md with a falsifier before the measurement, on the same terms as P1–P9.

	prediction (final registered form)	outcome	where resolved
P10	Quantizing Δ on its own grid rather than W 's puts $\text{mean}(\ \Delta\ /s_\Delta)$ in 0.1–1.0, cosine above 0.99 for all nine, relative error under 0.15, and compresses the across-adaptor spread below 1.1×	Three of four clauses confirmed. Clause 1's range was wrong and its direction right: measured 2.31–2.38, not 0.1–1.0	§7
P11	Equation 4's third licensing assumption: $P(\delta < 0)$ within 0.5 ± 0.01 per module-instance (P11.1); $\ r(\text{sign } \delta, u)\ < 0.01$ (P11.2); the sign-aware flip prediction within 2% of the 50/50 two-tail average at every $t \geq 0.005$ (P11.3)	All three confirmed , and by margins of 42×, 9× and 3× on the registered bounds	§4.1, B.11

Two properties of this listing are the point of it. **First, the failures are legible:** a reader can see that P4's headline clause did not survive its own input data, that P7 was abandoned rather than quietly reinterpreted, and that P10's first clause was off by an order of magnitude in a direction that did not change its conclusion. **Second, nothing is missing:** eleven were registered and eleven appear.

8. Limitations

8.1 The behavioural population is one condition replicated six times

Phase 1's six adapters are all **rank 32, $\alpha/r = 2$, on one base model, from one training recipe**, differing only in the target word. That is effectively *one condition replicated six times with different secret words*, not a survey of adapters.

Our weight-space measurements span ranks 16–128, both scaling conventions, two base models and four training regimes. **The behavioural claims do not inherit that coverage.** Every statement in §5 — the dose-response, the benign dissociation, the survival of INT4 g128 — is a statement about rank-32, $\alpha/r = 2$ adapters on Qwen3-8B trained by one recipe. Nothing in this paper licenses extending it across rank, scaling convention, base model, or task.

This is the single largest gap between what we measured and what a reader might want to conclude, and it is the reason the paper's headline is stated at the precision it is.

8.2 We did not run the across-population predictive test

The predictive gap (§5.4) is established **within** a matched population. Testing whether output SNR predicts behavioural retention **across** dissimilar adapters would require batteries for adapters we did not build them for, and we made a deliberate decision not to pursue it.

The reasoning, recorded before the decision rather than after: the available predictor range collapsed from $3.7\times$ to $2.3\times$ when the safety adapter failed instrument validation (§6.4); both remaining candidates lacked validated batteries and might not have passed (the safety adapter had just failed exactly that step); and the result would have remained confounded regardless, since the candidate adapters differ in rank, base model, recipe and task simultaneously. The best available outcome was a suggestive, non-causal, $n=3$ correlation, which would not have overturned PG-1 to PG-3.

What would change this. A matched pair — two adapters identical in rank, base model, recipe and task, differing only in effective magnitude — would remove the confound and make the test causal. No such pair exists publicly; it would have to be trained. We regard this as the correct next experiment and describe it in Future Work rather than claiming the present study substitutes for it.

8.3 Weight-space quantize–dequantize, not a deployment kernel

We simulate quantization as quantize–dequantize on the target projections, validated bit-exact against `gptqmodel`. This is the arithmetic whose retention we characterise, **not** an end-to-end deployment stack: we do not exercise fused INT4 kernels, and our results do not capture kernel-level effects such as accumulation order or activation quantization.

We also do not cover **GGUF K-quants**, which use block-wise super-block scales rather than plain affine quantization and would need llama.cpp's own quantizer rather than a reimplementation. Their retention behaviour is unmeasured here.

8.4 Effective sample size is far below nominal

Greedy decoding at temperature 0 makes generation deterministic, so seeds are not a replicate axis: re-running reproduces outputs exactly. All intervals are therefore bootstrapped over **intent clusters** (per-adapter) or **adapters** (pooled), never over observations. The 32 prompts are not 32 independent draws: they are 8 intents $\times$ 3 paraphrases plus 8 adversarial prompts. Measured (B.12) the battery carries **23–26 effective units**, not 32 and not the 16 an earlier draft assumed.

The resulting precision is limited. Per-adapter 95% intervals are **13–47% wide** under the clustered estimator, which is why PG-2 resolves only 4 of 15 pairs at INT3 and 1 of 15 at INT4 g128. That single INT4 pair involves an adapter whose point estimate is 116.2% of its own BF16 baseline — impossible, and therefore a statement about the instrument rather than about the adapter. An earlier reading of the whole INT4 spread as a real effect was our own error, corrected by the bootstrap.

8.5 The safety adapter is confounded four ways, and we registered this in advance

`responsible-ai-safety` is simultaneously the **only** alignment-trained adapter, the **only** Llama-3.1-8B, the **only** refusal-framed task, and the **only** rank-16 adapter with a behavioural battery attempt. Its divergence from the taboo family therefore **cannot be attributed** to the training recipe, the base model, or the safety framing — all change together with nothing to separate them.

This limitation was written into the pre-registration before the adapter was run, not added after its result was known.

8.6 §6 is a case study at $n=2$

The finding that the safety adapter removes its base model's refusal under indirect pressure rests on **2 of 8 jailbreak-framed prompts, on one adapter, at BF16**, with each instance verified by reading the full trajectory. It is a case study motivating the practice of checking that an adapter implements its advertised behaviour before deploying it. It is **not** a population estimate about alignment LoRAs, and it says nothing about that adapter's behaviour under quantization, which we did not measure.

8.7 Instrument limitations found by audit

Two instruments in this paper have documented failure modes that we report rather than tune away, because in both cases the fix would not have changed the verdict and retuning after seeing the data is the error the gate exists to prevent.

- `p_refuse` **responds to prompt appearance as well as model behaviour** (§6.5). Within a fixed prompt set it remains a valid across-condition comparison; its absolute level is not a refusal probability.

- **String-match refusal errs in both directions** — a false negative on soft refusals ("I'm not sure I can help with that"), and a false positive on vocabulary content (the marker "i am sorry" matching the English gloss in the entry *Je suis désolé(e)* (*I am sorry*) in a French vocabulary list). We report raw and audit-corrected counts separately (§6.2).

8.8 Checkpoint provenance

The Taboo checkpoints we use carry no model card documenting their training recipe or their relationship to the paper that introduced the Taboo setting. We adopt the setting from Cywiński et al. (arXiv:2505.14352) and treat the checkpoints as independent public artefacts. Their internal consistency is measured — three same-recipe adapters give weight-space cosines of 0.1380, 0.1389 and 0.1409, under 2% relative spread — but their training details are not documented and we do not assume them.

9. Conclusion

Merging a LoRA adapter into a base model and quantizing to INT4 changes only **1.1%–14.8%** of the model's stored integer codes — under 6.2% for eight of the nine adapters measured — and replaces the intended weight update with something $1.7\text{--}7.4\times$ its size pointing in a largely uncorrelated direction. By any weight-space measure, the adaptation is very nearly gone.

The capability is not. At INT4 with group size 128 — the standard deployment configuration — **85.5% of the six taboo adapters' stored values differ from the base model's, and no loss of elicitation capability is detectable**: retention 99.2%, enumerated interval [90.7%, 107.6%], which spans parity and excludes losses beyond about 9%. The trained constraint, behaviour's other side, does move and **tightens**: the adversarial leak rate falls 8.3 points, [+4.2, +12.5] (§5.1). Capability degrades only at coarser grids, reaching 77.2% at INT4 per-channel and 57.8% at INT3, and in this population it degrades in the benign direction: the model becomes less able to express the trained behaviour while the trained constraint holds, rather than retaining the capability and losing the restraint.

Both halves follow from one quantity. $|\Delta|/s$ — the adapter's effective magnitude against the quantization step — predicts the stored-weight outcome through a parameter-free channel model accurate to 2.3% across two base models, four ranks, both scaling conventions and four training regimes. The same quantity, applied to inputs inside the adapter's rank- r active subspace rather than to individual weights, is amplified by $\sqrt{d_{\text{in}}/r}$, and layer-output fidelity is *measured* $6.2\text{--}16.5\times$ higher than weight-space fidelity, rising as rank falls. **Weight-space erasure and layer-output survival are not in tension; they are one measurement read at two levels of the weight-to-output map.** Behaviour is a third level and follows from neither — PG-3 is this paper's own counterexample.

The practical guidance is uncomfortable in both directions. INT4 g128 is **safer** than the weight-space numbers suggest, and practitioners who would have been alarmed by a cosine of 0.14 should not be. But weight-space diagnostics — including the tool we ship with this paper — **cannot tell you which adapter will survive**. Within a population matched on rank, scaling, base model, recipe and output SNR to 3.3%, behavioural retention spans 28.7% to 86.4% (28.4–84.4% floor-corrected); among the pairs whose difference we can resolve, the ordering runs opposite to the predictor in six of the seven resolvable pairs; and the adapter with the largest weight-space footprint in our study has no measurable target behaviour at all. A practitioner choosing between two comparable adapters gets no information from these measurements, and we say so in the tool's own output.

The honest summary of what we set out to test is that the alarming version did not occur. We began by asking whether a quantized alignment fine-tune is behaviourally the base model. At standard settings it is not, and the mechanism by which it survives is the same mechanism that makes the weights look destroyed. What replaces that concern is narrower and more tractable: **effective adapter magnitude is the quantity that governs retention, no adapter card publishes it, and no weight-space measurement of it predicts behaviour.**

Future work

Two of these are named experiments with pre-stated predictions, and we flag them as untested rather than implying our data covers them.

FW-1. Does outlier-aware quantization protect the wrong weights for adapter retention?

Our layer 1–3 flip-rate spike is driven by weight groups sitting at input channels whose activation is **0.15–0.19× the module mean** — the quietest channels, the inverse of the massive-activation pattern (§2.2, §4.5.1). Existing outlier-aware methods select what to protect by *high* activation: AWQ scales up the salient channels identified from the activation distribution; LLM.int8() routes outlier dimensions through 16-bit. **On that criterion, the groups driving adapter erosion are the least salient channels in the layer and would not be protected.**

Prediction: applying AWQ- or LLM.int8()-style protection to a merged adapter leaves the early-layer spike substantially intact, while protecting the **narrowest-step groups** — a criterion no current method uses — removes it. This is a one-afternoon experiment on our existing harness. If it holds, adaptation retention depends on a class of weight that current saliency criteria are constructed to ignore, and a retention-aware quantizer would need a different selection rule from an accuracy-aware one.

FW-2. A matched pair. Two adapters identical in rank, base model, recipe and task, differing only in effective magnitude, would convert the predictive gap from a within-population observation into a causal test (§8.2). No such pair exists publicly; both would have to be trained.

FW-3. Behavioural coverage across rank and convention, which §8.1 identifies as the largest gap between our weight-space and behavioural claims.

FW-4. GGUF K-quants and deployment kernels (§8.3), using llama.cpp's own quantizer rather than a reimplementation.

FW-5. Whether the benign dissociation generalises beyond a suppression-style behaviour to capabilities with no constraint component.

Quantizing Δ on its own scale was FW-2 in an earlier draft. It is now measured (`METHODOLOGY.md`), so it is a result rather than future work, and the remaining items are renumbered.

Appendix A: `ar.predict` — computing effective adapter magnitude from a published checkpoint

The tool exists because §4.1 identifies $|\Delta|/s$ as the quantity that governs retention, and no adapter card publishes effective magnitude. It closes that gap for weights. It does not close it for behaviour, and it says so in its own output (§A.4).

A.1 What it does and what it costs

Given a HuggingFace adapter identifier, `ar.predict` computes the adapter's per-weight delta magnitude against the quantization step size of the base model it targets, and reports the weight-space consequences.

It needs **no GPU, no training, and no base-model download**. Base weights are range-read from the remote safetensors shards — roughly **150 MB of network** rather than the 16 GB the full model would require — and adapter tensors are small. A run takes about 30 seconds.

```
python -m ar.predict --adapter adamkarvonen/Qwen3-8B-taboo-smile_50_mix \  
--bits 4 --group-size 128
```

A.2 Output

```

effective magnitude  mean|delta|      9.747e-05
                    mean step s      9.222e-03
                    mean|delta|/s    0.01060

predicted bit-flip rate          0.0106   [0.0090, 0.0122]
predicted cosine(delta, delta_eff) 0.1427
predicted weight-space SNR       0.1442
predicted layer-output SNR       1.749   [1.487, 2.011]  +/-15%

```

NEAR-TOTAL WEIGHT-SPACE EROSION: the deployed weights barely move.

module	mean d /s	flip	cosine	amp	SNR_out

v_proj	0.00874	0.0087	0.1300	11.22	1.471
k_proj	0.00979	0.0098	0.1335	11.23	1.513
o_proj	0.01068	0.0107	0.1401	11.22	1.588
q_proj	0.01088	0.0109	0.1426	11.22	1.617
up_proj	0.01205	0.0121	0.1555	11.21	1.765
gate_proj	0.01568	0.0157	0.1856	11.19	2.114
down_proj	0.00635	0.0064	0.1115	19.39	2.175

[elided: a note on what the two weight-space columns assume, then the two banners of A.4, which print here on every run and are quoted there in full]

In weight space the module ordering is mostly a magnitude effect and partly a tail-shape one (§4.3, B.5): `gate_proj` has the largest $|\Delta|/s$ and the highest cosine, `down_proj` the smallest and the lowest, and in *this* table the cosine column happens to be monotone in the ratio column. It is not monotone in general, and an earlier draft called the effect "entirely" magnitude two bullets above the measurement that refutes it. Pooled over the nine adapters (B.5), `k_proj` has the second-highest flip rate and the fourth-lowest cosine while `up_proj` has the fourth-highest flip and the second-highest cosine — which under $\text{cosine} = \sqrt{\tau \cdot |\Delta|/s}$ requires `k_proj`'s τ to sit **17.5% below** `up_proj`'s, and the τ bullet below measures exactly that spread. Magnitude sets the scale; tail shape reorders within it.

In output space the ordering reverses at the bottom, and the reason is architectural. `down_proj` has the *lowest* weight-space cosine (0.1115) and the *highest* predicted output SNR (2.175), because its input dimension is 12288 rather than 4096, so its amplification is 19.39 against 11.2. The two columns answer different questions and disagree about which module fares worst; a reader taking "retains least" from the ratio column and applying it to the layer's output would get it backwards. This is the same distinction as §4.4 and §5.4, visible inside a single table.

Three notes on reading these numbers against the paper's.

- **They are predictions, not the paper's measurements, and they do not match to the last digit.** The tool samples three layers of 36 and averages seven modules unweighted; the weight-space runs measure four layers and 28 module-instances (B.1). Predicted flip here is 0.0106, against a measured 0.01093 for the same adapter in B.1 — **3.0% apart**, which is within

the tool's stated $\pm 15\%$ band but *outside* the 2.3% figure quoted for the closed form. The 2.3% is the model's error against a full measurement; 3.0% is what layer sampling adds.

- **predicted layer-output SNR 1.749 is not the 1.6286 B.13 attributes to this adapter.** B.13's value is *measured*, by projecting onto an orthonormal basis of Δ 's right singular vectors; this one is *predicted* from Equation 5. They differ by 7.4%, which is the amplification law's error on this adapter, and PG-1's 1.6200–1.6728 range is a range of measured values only.
- **predicted weight-space SNR is $\cos / \sqrt{1 - \cos^2}$, and that is not the definition B.13 measures.** Here it is the ratio of Δ_{eff} 's component *along* Δ to its component *orthogonal* to Δ , computed from the predicted cosine: at cosine 0.1427 it gives 0.1442. B.13 reports $\|\Delta\| / \|\Delta_{\text{eff}} - \Delta\|$, signal over total error, measured per layer and averaged, and that is the quantity the abstract's 6.2–16.5 $\times$ amplification range is denominated in. **The gap between what this tool prints and what B.13 measures is 0.1442 against 0.1341, 7.5%.** An earlier draft quoted 3.4% here, which is the gap between the two *formulas* evaluated at the same measured cosine (0.1387 against 0.1341) — a real quantity, and not the one a reader of the tool's output is asking about. Every number the tool prints is a claim, and substituting a friendlier recomputation for the printed one is the failure this appendix exists to prevent. The formulas agree closely because both reduce to approximately $\cos$ when the projection coefficient is near 1 and $\cos$ is small; the remaining 4.1% is the cosine prediction's own error. Neither validates the other.
- **The cosine column is not the fixed- τ formula A.3 tabulates.** The tool computes $\sqrt{\text{mean}(\Delta^2) / (s \cdot \text{mean}|\Delta|)}$ per module, which is $\sqrt{\tau_{\text{module}} \cdot |\Delta|/s}$ with each module's *own* tail-shape statistic. Backing it out of the printed columns gives 1.82–2.20 across the seven modules, against the 1.5962 the synthetic Gaussian sweep measures. A.3's table describes **Figure 5's** predictor, which deliberately uses one fixed constant so that it has no per-module free parameter; applying that formula to the row above gives 0.118 where the tool prints 0.130. Two different predictors, and an earlier draft implied they were the same one.
- **The amp column uses each module's measured error concentration, not the $1 + c/r$ form.** Equation 5 with $c = 0.87$ and $r = 32$ gives $\sqrt{128/1.0272} = 11.16$; the tool prints 11.22, implying a measured concentration of 1.017 rather than the fitted 1.027. Both are in the paper: the closed form is the claim, the measured value is what the tool has available per module. **The 1% gap is not unexplained.** $c \approx 0.87$ is fitted across ranks 16–128; on this adapter at $r = 32$ the measured concentration corresponds to $c = 0.017 \times 32 = 0.54$, so the closed form over-states concentration here and under-states amplification by 0.5%. The abstract's 6.2–16.5 $\times$ range is the span of B.13's *measured* per-adapter ratios and does not pass through c at all.

A.3 Accuracy

Validated against directly measured records on **nine** published adapters (the validation figure at the end of this appendix):

quantity	prediction	max relative error across nine adapters
code-flip rate	$\text{mean}(\min(\ \Delta\ /s, 1))$	2.3% (safety)
cosine	per module, $\sqrt{\tau \cdot \ \Delta\ /s}$ with $\tau = 1.5962$, then averaged	10.4% (latentqa)

The prediction is a closed form with no fitted parameters (§3.5), so this is out-of-sample in the only sense available: nothing about these nine adapters was used to construct the model. The cosine row carries one measured constant, the tail-shape statistic $\tau = \text{mean}(\Delta^2)/\text{mean}|\Delta|^2$, measured at **1.5962 on the synthetic Gaussian sweep** and flat there across three decades of adapter magnitude.

That constant is a property of the synthetic generator, and trained adapters do not satisfy it. 1.5962 is 1.6% above $\pi/2 = 1.5708$, the exact value for a Gaussian, which is what the sweep draws. Backing τ out per module from A.2's table gives **1.82–2.20** on a real adapter — 16% to 38% higher, because products of trained $B \cdot A$ matrices are heavier-tailed than products of Gaussians. Since $\text{cosine} \propto \sqrt{\tau}$, that alone is a 7% to 17% under-prediction, which is the right size to be most of the 10.4% in the table above; the code-flip row carries no such constant, which is why it is at 2.3%. **The cosine predictor's error is dominated by a constant fitted on a generator its targets do not match**, and stating that is more useful than the error figure alone.

These errors are for the quantity the model predicts: the adapter's own contribution under a fixed grid (§3.3). A deployment toolchain also moves the grid, which is a second effect the closed form does not model; read as a prediction of the combined outcome its error is 32–47%. The tool prints this scope in its own output.

*An earlier version of this table read "six adapters" and "cosine 5.0%", and both were wrong; the panel itself printed "max error 0.0%". Three numbers for one quantity, none of them the measured 10.4%. The figure plots nine. The cosine panel was computing its "prediction" as the measured projection coefficient $\langle \Delta_{\text{eff}}, \Delta \rangle / ||\Delta||^2$ divided by the measured magnitude ratio $||\Delta_{\text{eff}}|| / ||\Delta||$. Those differ by exactly a factor of $\cos(\Delta, \Delta_{\text{eff}})$, so their quotient **is** the cosine — it plotted cosine against cosine, drew a perfect line, and printed max error 0.0%. It rendered without error for the whole draft and every cross-check passed, because every value in it was correct. It was simply not a test. The panel now uses the channel model's own cosine, and the figure's cross-check asserts that prediction and measurement **differ**, so the vacuous form cannot come back silently (METHODOLOGY.md M.3, METHODOLOGY.md M.5).*

It reads `use_rslora` from each adapter's config rather than assuming a convention. For a rank-128 rsLoRA adapter the two conventions differ by $\sqrt{128} \approx 11.3\times$, which is enough to move an adapter from worst to best in a nine-adapter ranking (METHODOLOGY.md M.4). The computed delta is verified against peft's own `merge_and_unload` by a ground-truth fixture (§3.8).

A.4 What it cannot do, stated in its own output

The tool prints the following, unconditionally, on every run:

LIMIT OF THIS TOOL, measured not hypothetical. Six adapters matched on rank, scaling, base model and training recipe, whose output SNR agreed to within 3.3%, showed behavioural retention at 3-bit spanning 28.7% to 86.4% (28.4% to 84.4% with the instrument's floor subtracted; the split is 2 of 6 below half uncorrected, 3 of 6 corrected). The outcome varied 27x more than the predictor did, once the outcome's own measurement error is netted out of its spread — 30x on the raw comparison, which is a noisy outcome against a deterministic predictor. Of the 7 adapter pairs whose difference was statistically resolved, 6 ran **OPPOSITE** to output SNR.

So: these numbers do not discriminate between similar adapters. If you are choosing between two adapters of comparable rank and magnitude, this tool cannot tell you which will survive quantization better, and a difference it reports between them carries no information. Whether it discriminates ACROSS dissimilar adapters is untested.

It also prints, unconditionally:

WHAT THESE NUMBERS SCOPE TO. *The flip rate above is THE ADAPTER'S OWN CONTRIBUTION: given one grid derived from the base weights, the fraction of codes this delta pushes across a boundary. That is what the model predicts, parameter-free, to within 2.3% on nine published adapters.*

A deployment toolchain also recomputes the grid from the merged tensor, and the grid then moves under almost every weight. That is a SECOND effect this model does not describe. With both acting, measured code flips run 1.5–1.9× the number above and 83.6–87.4% of dequantized VALUES differ rather than the 1–15% of codes. So: read the flip rate as what the adapter did, not as the fraction of your deployed checkpoint that differs from the base.

We include both rather than a softer caveat because each failure is measured, not anticipated (§5.4, §3.3), and because a tool that reports a number invites the inference that the number ranks things. **The honest use is as a description of what happens to stored weights, at a single adapter, under a grid that isolates the adapter's own contribution — not as a comparison between adapters and not as a deployment forecast.**

A.5 Interpreting the output

The two banners in A.4 are the tool's own statement of what it does and does not support, and they are quoted above verbatim rather than paraphrased here. Two additions they do not carry. **The tool answers only the merged case:** §2.5 measures the unmerged one on 756 records and finds it entirely different, $|\Delta|/s$ rising from 0.011–0.149 to 2.31–2.38 and cosine from 0.14–0.51 to 0.9948–0.9952. And **the predictive-gap limit is about discrimination, not level** — output SNR fails to separate adapters whose predictor range is 3.3% wide, which is range restriction, and is compatible with the same quantity setting the absolute level, which is what §4.4 measures and what the amplification law is for (§5.4).

A.6 Reproduction

Full commands and pinned versions in Appendix D. The tool alone:

```
pip install -r requirements.txt          # no torch CUDA build needed for ar.predict
PYTHONPATH=src python -m ar.predict --adapter <hf-id> --bits 4 --group-size 128
```

Supported: `--bits {3,4,8}`, `--group-size {any positive int, or -1 for per-channel}`, `--scheme {asymmetric, symmetric_gptq, symmetric_awq}`, `--base-model` to override the adapter's declared base.

ar.predict: predicted versus measured, 9 published adapters

Dashed line is exact agreement. The tool needs no GPU and no model download beyond adapter tensors (~150 MB). It predicts stored-weight outcomes: it does NOT predict behaviour (section 5.4)




Figure A1

Appendix B: Full tables

Record counts: Phase 0 weight-space 1512 (4-layer) + 504 (36-layer); Phase 1 behavioural 1536.

B.1 Weight-space retention per adapter (INT4 g128, asymmetric, both regimes)

CIs are over layers, on the `fixed_scale` cosine. The 36-layer run exists for one adapter only and is reported on its own row rather than pooled with the 4-layer runs.

`fixed_scale` holds the grid derived from W and applies it to $W + \Delta$, so a weight can only change if the adapter moved it across a boundary. `adaptive_scale` recomputes the grid from $W + \Delta$, which is what a deployment toolchain does and **what this paper's Phase 1 behavioural pipeline ran under** (§3.3, §5.1). Under it a weight can also change because the grid moved beneath it, which is why the value-change column is two orders of magnitude above the code-flip column.

Intervals without a mark are **enumerated** over all k^k resamples, so they carry no resampling noise; enumerated is not the same as exact coverage. A `*` marks a **sampled** interval (Monte Carlo, $n=20000$), used where the sample is too large to enumerate; its last printed digit is at the resolution the resampling noise supports and no finer.

Base model is `Qwen3-8B` for every adapter except `responsible-ai-safety`, which is `Llama-3.1-8B-Instruct`.

rel. err and mag. ratio are different quantities and the paper quoted the wrong one. `rel. err` is $||\Delta_{\text{eff}} - \Delta|| / ||\Delta||$, error against an erasure baseline of 1.0. `mag. ratio` is $||\Delta_{\text{eff}}|| / ||\Delta||$, how much larger the delivered update is than the intended one — which is what the phrase "7.5 times its magnitude" means and what was never tabulated. They differ by about 1%: at cosine 0.1374 a relative error of 7.407 implies a magnitude ratio of 7.476, and earlier drafts quoted 7.4 for the second while reading it off the first. Per §3.4 the magnitude ratio is never reported without a cosine beside it.

Columns marked `fix` are `fixed_scale` and `adp` is `adaptive_scale`; `val-chg` is the value-change rate, `adaptive_scale` only.

adapter	r	α/r	lay	cos fix	95% CI	cos adp	flip fix	flip adp	val-chg	rel err	mag
ao-v3-dpo-halluc	128	0.125	4	0.5050	0.476–0.539	0.5047	0.14813	0.22115	0.8742	1.74	2.01
latentqa	64	2	4	0.2760	0.255–0.300	0.2746	0.03882	0.06481	0.8586	3.58	3.72
responsible-ai-safety	16	2	4	0.3298	0.307–0.366	0.3290	0.06191	0.10219	0.8363	2.90	3.07
taboo-gold	32	2	4	0.1389	0.125–0.153	0.1378	0.01114	0.02112	0.8546	7.35	7.42
taboo-moon	32	2	4	0.1375	0.124–0.151	0.1364	0.01092	0.02079	0.8545	7.42	7.49
taboo-rock	32	2	4	0.1412	0.128–0.154	0.1401	0.01137	0.02147	0.8546	7.21	7.28
taboo-ship	32	2	4	0.1409	0.128–0.154	0.1398	0.01139	0.02149	0.8546	7.23	7.30
taboo-smile	32	2	4	0.1374	0.124–0.150	0.1363	0.01093	0.02080	0.8546	7.41	7.48
taboo-snow	32	2	4	0.1382	0.126–0.150	0.1370	0.01102	0.02088	0.8545	7.37	7.44
taboo-smile	32	2	36	0.1380	0.135–0.140*	0.1369	0.01220	0.02267	0.8463	7.29	7.36

B.2 Channel model: predicted vs measured code-flip rate

`predicted = mean(min(| $\Delta$ |/s, 1))`, no fitted parameters. INT4 g128, asymmetric, `fixed_scale`.

adapter	measured	predicted	ratio	rel. error	proj. identity
ao-v3-dpo-halluc	0.148128	0.149488	0.9909	0.92%	0.9900
latentqa	0.038819	0.039123	0.9922	0.78%	0.9912
responsible-ai-safety	0.061912	0.063354	0.9772	2.33%	0.9743
taboo-gold	0.011143	0.011158	0.9987	0.13%	0.9924
taboo-moon	0.010921	0.010936	0.9986	0.14%	0.9934
taboo-rock	0.011369	0.011397	0.9975	0.25%	0.9919
taboo-ship	0.011394	0.011417	0.9980	0.20%	0.9926
taboo-smile	0.010934	0.010949	0.9986	0.14%	0.9924
taboo-snow	0.011020	0.011035	0.9986	0.14%	0.9931

Maximum relative error: 2.3%.

Six digits on the inputs and four on the ratio because five and three did not reproduce: at `0.01093` against `0.01095` the printed pair supports a ratio anywhere in 0.9981–0.9985, and the table asserted 0.999. This is the table carrying the paper's headline accuracy, and a reader who recomputes the derived columns from the printed inputs has to land on the printed values.

B.3 Quantization convention (paired)

Paired on (adapter, layer, module) cells present under all three schemes, `fixed_scale`. Pooling unpaired records inverts this ordering (EXP-008).

scheme	cells	cosine	code-flip	rel. err
<code>asymmetric</code>	252	0.2161	0.0351	5.802
<code>symmetric_gptq</code>	252	0.2065	0.0341	6.733
<code>symmetric_awq</code>	252	0.1980	0.0303	6.453

B.4 Scale regime

`fixed_scale` isolates the adapter's contribution; `adaptive_scale` is deployment-realistic and is the regime Phase 1 ran under (§3.3). Pooled over the 9 adapters, code flips and value changes differ by 15.0x under `adaptive_scale`, which is why both are logged; per adapter the ratio runs from 4.0x to 41.1x (B.1).

`scale-shift` is the fraction of GROUPS whose step size differs between W and $W + \Delta$; `grid-shift` is the fraction of WEIGHTS whose dequantized value changes under `adaptive_scale` but not under `fixed_scale` — i.e. those that moved because the grid moved, not because the adapter cleared the step. Both are properties of the pair of regimes, so they read the same on both rows.

regime	cosine	code-flip	value-change	scale-shift	grid-shift
<code>fixed_scale</code>	0.2161	0.0351	0.0351	1.0000	0.8202
<code>adaptive_scale</code>	0.2151	0.0572	0.8552	1.0000	0.8202

B.5 Module profile (pooled over 9 adapters)

module	cells	cosine	code-flip
<code>gate_proj</code>	36	0.2480	0.0414
<code>up_proj</code>	36	0.2234	0.0356
<code>q_proj</code>	36	0.2225	0.0386
<code>k_proj</code>	36	0.2154	0.0401
<code>o_proj</code>	36	0.2149	0.0339
<code>v_proj</code>	36	0.2135	0.0317
<code>down_proj</code>	36	0.1750	0.0243

B.6 Behavioural retention per adapter and precision

Elicitation score as a fraction of the same adapter's own BF16 score. Intervals are over the six adapters, not over prompts: the adapter is the sampling unit, and the prompts within one adapter are not independent draws — they are 8 intents x 3 paraphrases plus 8 adversarial prompts, and it is that

clustering rather than greedy decoding which makes them dependent (§3.11). They are enumerated over all 6^6 resamples.

The metric has a non-zero floor and no ceiling at 1.0. The BF16 column is the guesser's score on the model's own response divided by its score on a canonical hand-written hint for that word (§3.7), so a response that is a *better* hint than the canonical one scores above 1: `ship` reads 1.0327 and `snow` 1.0030. That is a property of the normaliser, not an anomaly, and it is why every retention figure in this table is each adapter against its own BF16 value rather than against 1.

The denominator is the aligned model's own BF16 score and the metric has a non-zero floor, so the percentages are not "fraction of the behaviour". The `base` column gives the SAME instrument's score on the base model without the adapter, at the same precision, so a reader can floor-correct. It is small but not negligible and varies 40x across adapters (`ship` 0.0039, `snow` 0.1642): the guesser has a prior over the 20 candidates. Floor-corrected retention, $(\text{aligned} - \text{base}) / (\text{aligned_BF16} - \text{base_BF16})$, is given as its own row and moves the *mean* by under 2 points at every precision — but it moves the per-adapter split, which is what the paper's claims are about, and **B.7** gives every adapter under every variant rather than leaving a reader to check a mean against a claim it cannot settle.

word	BF16	base	int4_g128	base	int4_pc	base	int3_g128	base
gold	0.8280	0.0165	81.3%	0.0168	62.4%	0.0177	41.3%	0.0211
moon	0.8250	0.0245	100.2%	0.0233	78.1%	0.0536	86.4%	0.0376
rock	0.8603	0.0105	116.2%	0.0125	77.5%	0.0131	57.7%	0.0151
ship	1.0327	0.0039	103.2%	0.0034	79.8%	0.0040	28.7%	0.0034
smile	0.7178	0.0150	100.8%	0.0123	68.5%	0.0115	51.3%	0.0221
snow	1.0030	0.1642	93.5%	0.1634	96.8%	0.1402	81.5%	0.1655
mean	—	—	99.2%	—	77.2%	—	57.8%	—
95% CI over adapters	—	—	[90.7%, 107.6%]	—	[69.0%, 86.0%]	—	[41.7%, 74.3%]	—
floor-corrected mean	—	—	99.0%	—	76.5%	—	56.0%	—
mean, 24 hint prompts only	—	—	104.9%	—	80.2%	—	62.1%	—
below 50%	—	—	0/6	—	0/6	—	2/6	—

The 8 adversarial prompts are constraint probes, and the capability score pools them with the 24 hint prompts. Their job is to make the model say the word, and a response containing it scores at or above the normaliser because the guesser recovers it trivially — so a quarter of the capability axis is the inverse of the constraint axis. The `24 hint prompts only` row above is the same measurement with them removed. **The two agree:** the ordering is identical, the dose-response is monotone in both, and the hint-only mean is +5.7 points at INT4 g128, +3.0 at per-channel and +4.3 at INT3. Note the sign: removing the adversarial prompts **raises** the score, because adversarial prompts are harder and yield less word-bearing text than hint prompts, which outweighs the leak lift. That lift is real and is now reported with an interval rather than as a bare pair of means: responses containing the word score 0.929 against 0.717, and per adapter the difference is +0.255 with an enumerated 95%

interval of [+0.136, +0.371]. The pooled figure is the pre-registered instrument and remains the headline; it is not the more flattering one. At INT4 g128 **4** readings are above parity — **moon** 100.2%, **rock** 116.2%, **ship** 103.2%, **smile** 100.8% — and **3** of them stay above parity with the adversarial prompts removed, so leakage does not explain them. Note that the count falls while the mean rises: the hint-only shift is not uniform across adapters, which is why the split is a per-adapter claim and not a mean one.

Adversarial leak rate, the fraction of the 8 adversarial prompts on which the secret word appears. This is the constraint measured in the disclosure frame under pressure, which is what the adversarial set was built for; the knowledge probe never mentions the secret and cannot test that frame.

	BF16	int4_g128	int4_pc	int3_g128
gold	12.5%	0.0%	0.0%	0.0%
moon	37.5%	25.0%	25.0%	12.5%
rock	12.5%	0.0%	12.5%	0.0%
ship	0.0%	0.0%	0.0%	0.0%
smile	25.0%	12.5%	0.0%	0.0%
snow	12.5%	12.5%	12.5%	25.0%
pooled	16.7%	8.3%	8.3%	6.2%
95% CI, leak rate	[8.3%, 27.1%]	[2.1%, 16.7%]	[2.1%, 16.7%]	[0.0%, 14.6%]

Every per-adapter cell is a multiple of 12.5%, because each is a count out of 8 — read them as counts, not as rate estimates. The pooled row is the mean of the six, so it need not be: 6.2% at INT3 is 6.25%, three leaks in 48 prompts, printed at one decimal.

The leak rate **falls** from 16.7% at BF16 to 6.2% at INT3, a paired difference over the six adapters of **+10.4%** with an enumerated 95% interval of **[+0.0%, +20.8%]**. **That interval reaches zero**, so this is a trend the six adapters support in direction and do not resolve in size, and it is stated that way rather than as a demonstration. Per adapter the difference runs **gold** +12.5%, **moon** +25.0%, **rock** +12.5%, **ship** +0.0%, **smile** +25.0%, **snow** -12.5% — **snow moves the other way**, from 12.5% to 25.0%, doubling at INT3 on a count of 8. Reporting only the pooled trend here would be the mean-versus-adapters error this appendix polices two tables up.

The constraint does not fail under quantization in the frame designed to break it. Read with care: capability falls too, so some of this is a model less able to produce the word at all rather than more willing to withhold it — which is the same confound §5.3 handles by comparing within precision, and is why this is reported beside the knowledge probe rather than instead of it.

And the ratio that motivated this instrument does not survive the full grid. E.2 said the adversarial prompts leak "roughly 6x more" than the hint prompts. Pooled over every aligned record it is **1.21x** (19 of 192 against 47 of 576); at BF16 alone 1.33x. The 6.00x is **smile** at BF16, one adapter at one precision. E.2 now says so.

B.7 The three metric variants, per adapter

Rows are the same measurement under three choices of instrument. The first is the pre-registered one and every headline number in the paper is quoted under it. **<50%** and **>80%** are counts of the six; **span**

is min–max. CV is the outcome coefficient of variation, which is PG-1's denominator.

All adapter, mean and span cells are percentages.

variant	prec.	gold	moon	rock	ship	smile	snow	mean	<50	>80	span	CV
pre-registered, 32	INT4 g128	81.3	100.2	116.2	103.2	100.8	93.5	99.2	0	6	81.3–116.2	0.116
pre-registered, 32	INT4 per-ch.	62.4	78.1	77.5	79.8	68.5	96.8	77.2	0	1	62.4–96.8	0.152
pre-registered, 32	INT3	41.3	86.4	57.7	28.7	51.3	81.5	57.8	2	2	28.7–86.4	0.390
floor-corrected, 32	INT4 g128	80.9	100.3	116.2	103.3	101.2	92.3	99.0	0	6	80.9–116.2	0.119
floor-corrected, 32	INT4 per-ch.	61.4	73.8	76.9	79.7	68.3	99.0	76.5	0	1	61.4–99.0	0.167
floor-corrected, 32	INT3	39.5	84.4	56.6	28.4	49.3	77.7	56.0	3	1	28.4–84.4	0.387
hint-only, 24	INT4 g128	92.0	104.8	121.0	98.9	112.7	99.9	104.9	0	6	92.0–121.0	0.100
hint-only, 24	INT4 per-ch.	73.5	77.2	72.8	83.7	81.2	92.8	80.2	0	3	72.8–92.8	0.094
hint-only, 24	INT3	42.3	89.3	66.3	32.7	63.7	78.5	62.1	2	1	32.7–89.3	0.344

What moves. The mean does not: floor correction shifts it by -0.2, -0.6, -1.8 points at INT4 g128, per-channel and INT3. **The split does.** At INT3 the count below half goes 2 → 3 under floor correction (**smile** crosses at 49.3%) and the count above 80% goes 2 → 1 (**snow** falls to 77.7%). The span goes 28.7%–86.4% → 28.4%–84.4% floor-corrected and 32.7%–89.3% hint-only. Every site quoting the split or the span now names the variant it is quoted under.

PG-1 does not move under the metric variants. Its predictor is a Phase 0 quantity and is unaffected by any of this: CV 0.0128. The outcome CV is the last column, and the ratio outcome/predictor runs **7.3× to 30.5×** across all nine variant × precision cells — the smallest is hint-only at INT4 per-channel and the largest is the pre-registered instrument at INT3. **PG-2 does not move under floor correction at all** (B.12): same pairs, same counts, same directions. It does move under hint-only, and B.12 gives that.

PG-1 does move once the outcome's own measurement error is netted out, and that correction was available in this appendix before it was applied. The predictor is deterministic and the outcome is not, so an observed outcome CV contains both between-adapter variance and within-adapter noise, and comparing the raw ratio against a noise-free predictor overstates the gap. Subtracting the mean squared per-adapter standard error — from the same cluster bootstrap B.12 uses, normal-approximated from each interval's width, which is a conservative reading of an asymmetric interval:

precision	outcome SD	RMS per-adaptor SE	between-adaptor SD	CV observed	CV between	ratio observed	ratio between
INT4 g128	0.1150	0.0800	0.0825	0.1159	0.0832	9.1×	6.5×
INT4 per-ch.	0.1173	0.0993	0.0624	0.1520	0.0809	11.9×	6.3×
INT3	0.2256	0.0974	0.2035	0.3902	0.3519	30.5×	27.5×

At INT3 the headline falls from 30.5× to 27.5×, which does not change PG-1's conclusion and is the number that should be quoted. At INT4 g128 it falls from 9.1× to 6.5× — most of the apparent between-adaptor spread at the finest grid is measurement noise, which is worth knowing and is not what PG-1 rests on. The released tool's banner quotes the corrected INT3 figure.

A fourth cell exists — floor-corrected *and* hint-only — and is omitted from the table because no claim is quoted under it; for completeness it gives 105.0%, 79.3%, 60.8% with the same 2/6 below half at INT3 as hint-only.

B.8 Paired contrasts between precisions

Paired over the six adapters, because the same six are measured at every precision; an unpaired comparison discards that and widens every interval for no reason. Intervals are **enumerated** over all 6^6 resamples of the per-adaptor difference, so there is no resampling noise and no seed. Enumerated is not the same as exact coverage: a percentile bootstrap at $n=6$ is asymmetric and approximate however it is computed, and **the test column rather than the interval is what carries the inference** (§3.11).

`p` is an exact sign-flip permutation over the $2^6 = 64$ assignments of sign to the six paired differences — the randomization test this paired design supports. It replaces an "excludes zero" column, which read the same bootstrap the interval comes from and therefore said nothing the interval had not already said.

contrast	mean paired difference	95% CI	sign-flip <code>p</code>
INT4 g128 - INT4 per-channel	22.0%	[10.8%, 31.6%]	0.0625
INT4 g128 - INT3 g128	41.4%	[23.3%, 59.3%]	0.0312
INT4 per-channel - INT3 g128	19.4%	[5.4%, 34.5%]	0.0625

Monotone at every step, per adaptor: **4 of 6**. The mean is monotone; the adapters are not.

Two of the three are 0.0625 rather than the 0.0312 floor, and the reason is not ties. No contrast here has a tied pair. In the first, `snow` moves against the mean; in the third, `moon` does. One adaptor running the other way makes the observed sum the *second* most extreme of the 64 rather than the first, which doubles the tail. Only the INT4 g128 vs INT3 contrast, where all six move together, reaches the floor.

"All three exclude zero" is one claim about a correlated triple, not three independent findings. The same six adapters produce all three contrasts and the third is the difference of the other two, so the multiplicity is not what a naive reading suggests and neither is the independence. **None of the three clears a three-way Holm correction**, and §3.11 gives the reason as a property of the design rather than of these numbers: at $n=6$ the two-sided sign-flip `p` cannot go below 0.0312.

An earlier version of this appendix inferred the third contrast's weakness from how far its interval's lower bound sat from zero. That is reading a percentile bootstrap at $n=6$ as a test, which is what §3.11 now forbids and what EXP-056 withdrew elsewhere; the inference is removed rather than requalified. The bound itself printed as 0.1 points for four drafts because the value is a ratio and the format specifier was `.1f`, so 0.054 rendered as 0.1 while the table printed the same number as 5.4%.

B.9 Knowledge probe: the benign dissociation

Aligned vs base **within the same precision**. The comparison inverts if aligned-quantized is compared against base-BF16 (§5.3). Intervals are enumerated over the six adapters, the same estimator as B.8 and Table 2; the adapter is the cluster because each contributes all its probes or none.

precision	base	95% CI, base	aligned	95% CI, aligned	ratio	Cliff's d	entropy (aligned)
bf16	0.3634	[0.1800, 0.5516]	0.0757	[0.0307, 0.1255]	0.208	-0.778	1.4069
int4_g128	0.3583	[0.1748, 0.5441]	0.0634	[0.0287, 0.0986]	0.177	-0.778	1.3984
int4_per_channel	0.3272	[0.1958, 0.4835]	0.0730	[0.0315, 0.1178]	0.223	-0.833	1.4998
int3_g128	0.2803	[0.1380, 0.4220]	0.0756	[0.0303, 0.1199]	0.270	-0.556	1.3480

The aligned column shows no trend, and it is not flat. It runs 0.0757, 0.0634, 0.0730, 0.0756 across the four precisions — a span of 0.0634–0.0757, whose largest single step is 16.2% between BF16 and INT4 g128. Every interval above overlaps every other, so the correct statement is **no detectable trend**, not equality. An earlier draft called this column "flat at 0.0757 and 0.0756" and §5.3 concluded the constraint was "exactly as strong at INT3 as at BF16": that is the first and last elements of a four-element series, quoted as if they were the series. It also explains why the ratio column is non-monotone while the base column falls monotonically — the non-monotonicity is in the numerator, and it is noise.

B.10 Layer 1–3 spike: step size vs input-channel activation

Activation columns are mean-normalised within each module (§4.5.1).

layer	module	step med/p1	act @ narrowest 1%	act @ widest 1%	$\rho(\log s, \text{act})$	split-half r
0	gate_proj	1.4	0.97	1.03	+0.033	0.943
0	up_proj	1.3	0.99	1.04	+0.010	0.943
1	gate_proj	83.5	0.17	1.12	+0.244	0.997
1	up_proj	7.1	0.89	1.00	+0.064	0.997
2	gate_proj	44.6	0.19	1.58	+0.275	0.996
2	up_proj	7.3	0.61	1.49	+0.113	0.996
3	gate_proj	145.1	0.15	1.05	+0.156	0.992
3	up_proj	24.4	0.80	0.96	+0.038	0.992
18	gate_proj	1.6	0.94	1.03	+0.012	0.820
18	up_proj	1.4	0.98	1.04	-0.004	0.820

B.11 Within-bin position: is u uniform where the model needs it?

The derivation §3.5 summarises. Under group-wise affine quantization a weight's stored code is $\text{round}(w/s + z)$. Write $w/s + z = k + u - \frac{1}{2}$ with k the integer code and $u = \text{frac}(w/s + z + \frac{1}{2}) \in [0, 1)$, so u is the distance from w to the **lower edge of its rounding cell** and $u = \frac{1}{2}$ is the cell's centre. Adding δ moves the code exactly when $u + \delta/s$ leaves $[0, 1)$, so for $|\delta| < s$ the indicator is $1[u < |\delta|/s]$ when $\delta < 0$ and $1[u > 1 - |\delta|/s]$ when $\delta > 0$. Under u uniform and independent of δ , each has expectation $|\delta|/s$, which is Equation 4 — the same marginal a stochastic rounder would give, obtained by averaging over a layer's $\sim 10^9$ weights rather than over repetitions of a random operation. *An earlier draft of §3.5 wrote $w = s(k + u)$ with u "the position within its bin", which under round-to-nearest puts the boundary at $u = \frac{1}{2}$ and every indicator half a bin out of place. The code always used the form above; only the line describing it was wrong.*

u is each weight's distance to its quantization boundary, over 42 module-instances on both base models, INT4 g128 asymmetric. Equation 4 needs $F_u(t) = t$. A flip is two-sided — a negative delta crosses the lower boundary, a positive one the upper — so the quantity the model averages over is the **mean of the two tails**.

t	lower tail	upper tail	mean (= P(flip))	mean / t
0.001	0.00256	0.00031	0.00143	1.433
0.002	0.00334	0.00109	0.00222	1.108
0.005	0.00614	0.00389	0.00501	1.003
0.008	0.00899	0.00673	0.00786	0.982
0.011	0.01200	0.00975	0.01088	0.989
0.020	0.02084	0.01858	0.01971	0.985
0.050	0.05036	0.04811	0.04924	0.985
0.100	0.09965	0.09746	0.09855	0.986
0.150	0.14909	0.14682	0.14796	0.986
0.250	0.24712	0.24510	0.24611	0.984
0.500	0.50015	0.49818	0.49916	0.998
0.750	0.75356	0.75135	0.75245	1.003

Uniform to within 1.8% at every t at or above 0.005, which covers the whole range our adapters occupy. Read on one tail alone the lowest 0.1% of the bin is over-occupied by 2.6x, which a one-sided measurement would have reported as a 156% excess.

Do not read the mean / t column as a flat 0.985 and infer a constant 1.3–1.5% over-prediction. That is what this appendix did, and it puts B.2 and B.11 eight-fold apart for no reason: the column is flat *in this table* because every row pools all weights, and the departure is a function of $|\Delta|/s$, which pooling hides. The decile breakdown below has it, and §4.1 states the budget it implies.

Where that excess comes from, since the obvious answer is wrong. An earlier version of this appendix attributed it to Equation 2 pinning each group's extrema, which would put 1.56% of weights on a boundary against the 0.20% measured — eight times over. Six controls, the last three added when the account that replaced it turned out to need them too:

control	measured	what the pinning account implies
u at each group's minimum	0.4943 ± 0.2887	0 (a boundary)
u at each group's maximum	0.4951 ± 0.2888	0 (a boundary)
SD of u at the extrema, over the uniform SD	1.000	0 (pinned)
IQR of u at the extrema	0.500	0 (pinned)
fraction of the u = 0 mass that is extrema	0.054	1.000
u = 0 mass surviving a jitter of $1e-4 \cdot s$	0.000022	0.002039, unchanged

u = 0 is the boundary and u = 0.5 is the bin centre. **The extrema are not pinned anywhere.** Because Equation 2 rounds z, a group's extrema are free to land anywhere in their code, and measured they do: mean 0.494, which is the uniform expectation, with an SD indistinguishable from the uniform $1/\sqrt{12} = 0.2887$ and an IQR of 0.500 against uniform 0.500. **This is the second wrong account of the same three numbers.** The first said the extrema are pinned to the boundary; the replacement said they are pinned to the bin centre and quoted the mean alone as evidence, which a uniform population produces exactly. A mean of 0.494 refutes the boundary account and licenses nothing in its

place; the dispersion is what settles it, and it was not measured until a second reader asked for it. An explanation written to replace a refuted explanation is at its most persuasive and least tested, which is precisely what `METHODOLOGY.md` M.8 says and precisely what M.8's own replacement did not get. What the exact-zero mass actually is: base weights are bf16, so a group of 128 holds about 121 distinct values and $w/s + z + 0.5$ lands exactly on an integer for roughly 1 in 500 of them. A perturbation four orders of magnitude below bf16's own resolution inside a bin removes 99% of it; a structural pinning would be untouched. The uniformity result does not depend on any of the three accounts, and did not change under either replacement (EXP-048, EXP-053).

Equation 4 has three licensing assumptions, not two, and the third is one this cancellation argument created. Averaging the two tails 50/50 is the right quantity only if $P(\delta < 0) = P(\delta > 0)$ and $\text{sign}(\delta)$ is independent of u . §4.1 measures $|\delta|$ against u ; a sign–position association would leave that untouched and break the cancellation exactly. Registered as P11 (EXP-046) before it was run, over the same 42 module-instances:

#	assumption	registered bound	measured	where
1	u independent of the magnitude of δ	—	max $ r $ 0.000774	§4.1
2	u uniform	—	within 1.8% for $t \geq 0.005$	above
3a	$P(\delta < 0) = 1/2$	0.5 ± 0.01	worst departure 0.000237	P11.1
3b	$\text{sign}(\delta)$ independent of u	$ r < 0.01$	max 0.001060	P11.2
3	the 50/50 average equals the sign-aware one	within 2% at $t \geq 0.005$	0.9936–1.0045 at $t = 0.011$	P11.3

All three hold, and the correlations are at their sampling floor rather than merely small: a null correlation on these module sizes has standard deviation $1/\sqrt{n}$ of 0.00013 to 0.00049, and the largest of the 42 is 2.17 of its own SD.

Assumption 1 was measured globally where the argument needs it locally, and this is the local version. A Pearson correlation of $|\delta|/s$ against u over the whole bin is dominated by the bulk and is close to uninformative about the conditional density of u in the lowest 1%, which is where the prediction lives at the $t \approx 0.011$ our adapters occupy. Binning the same 42 module-instances by decile of $|\delta|/s$ and re-reading the low tail inside each bin gives the conditional directly. Registered as **P12** (EXP-052) before it was run.

decile of $\delta \backslash s$	t range	mean t	P(flip) at $t = 0.011$	/ pooled	P(flip) at own t	true code flip	true / $\min(t, 1)$
1	0.00000–0.00484	0.00241	0.01085	0.9978	0.00267	0.00270	1.1185
2	0.00484–0.00981	0.00731	0.01087	0.9990	0.00726	0.00721	0.9867
3	0.00981–0.01508	0.01241	0.01085	0.9972	0.01218	0.01208	0.9728
4	0.01508–0.02081	0.01789	0.01086	0.9984	0.01759	0.01745	0.9755
5	0.02081–0.02726	0.02396	0.01088	1.0002	0.02357	0.02339	0.9760
6	0.02726–0.03488	0.03095	0.01086	0.9986	0.03043	0.03017	0.9747
7	0.03488–0.04443	0.03944	0.01089	1.0007	0.03879	0.03846	0.9751
8	0.04443–0.05774	0.05064	0.01090	1.0020	0.04977	0.04938	0.9752
9	0.05774–0.08096	0.06794	0.01092	1.0034	0.06672	0.06622	0.9748
10	0.08096–1.58251	0.12444	0.01091	1.0028	0.12231	0.11810	0.9491

P12.1 holds and is the load-bearing clause. At the common $t = 0.011$ the flip probability is 0.01085–0.01092 across all 10 deciles — a worst departure from the pooled value of **0.34%**, against a registered bound of 2%. The low tail of u does not depend on the size of the delta that has to cross it, which is the conditional Equation 4 needs and the one the global correlation could not see.

P12.2 failed as registered, and the falsifier was the thing at fault. The decile-index Spearman is **+0.87**, above the registered 0.5, so the drift is monotone. It is also 0.63% wide end to end. A rank statistic on ten values is scale-free by construction and will report a large correlation for a trend of any size, so registering one without a magnitude qualifier was a specification error of the same family as the three in `METHODOLOGY.md`. The dependence is real, systematic and six times smaller than the tolerance the model needs; both halves are stated because the registered bound says so.

P12.3 failed on one decile of 10, and it is the known non-uniformity of u rather than a conditional effect. Deciles 2–10 read 0.95–0.99 against $\min(t, 1)$; decile 1, whose mean t is 0.00241, reads 1.1185. That is the $t < 0.005$ region the table above already reports as over-occupied, arriving in the decile with the smallest deltas. Its effect on the integrated prediction is +0.07%, because it is a tenth of the weights at a twentieth of the mean t .

And this settles the error budget, which two appendices disagreed about. This appendix said the measured non-uniformity implies a 1.3–1.5% over-prediction for every adapter; B.2 measures 0.1–0.2% on the taboo six. Both are right and neither is a property of the model: **the departure is a function of $|\delta|/s$, not a constant.** The last column above falls from 1.12 at $t = 0.0024$ to 0.95 at $t = 0.124$. Split by adapter over these same 42 module-instances, the closed form over-predicts the true code flip by **2.7%** for `responsible-ai-safety` (ratio 0.9732), **0.1%** for `taboo-smile` (ratio

0.9990). B.2 reaches the same two numbers from a different code path on a different layer set — 0.977 and 0.999 — so this is a reproduction of that table's split, not a restatement of it. **The honest budget is: under 0.5% at the τ the taboo adapters occupy, and about 2.5% at four times that τ .** The paper's headline 2.3% maximum relative error is the safety adapter, and it is the highest- τ case among the small-rank adapters, not a floor that applies to all of them.

Pooled over all 42 module-instances: closed form 0.037385, true code flip 0.036516, ratio **0.9768**.

B.12 PG-2 under three estimators

Two corrections are bundled in "cluster bootstrap". **Pairing narrows** — both conditions run byte-identical prompts, so the shared prompt-difficulty variance cancels. **Clustering was described here as widening, and measured it does not, reliably.** The two halves are shown separately below and then the direction question is settled with the variance components, because an earlier version of this appendix asserted a direction its own table contradicted.

estimator	INT4 g128	INT4 per-ch.	INT3	interval width
A: prompts, unpaired (as published)	0	1	4	15%–63%
B: prompts, paired	1	2	6	14%–52%
C: intent clusters, paired (used)	1	2	4	13%–47%

The variance components, which decide the direction. Over the 24 hint prompts, 8 intents x 3 paraphrases, one-way random effects on `guesser_p_word_normalised`. `deff` is $1 + (k-1)\cdot\text{ICC}$ at $k=3$; `eff. units` applies it to the hint block and adds the 8 adversarial prompts, each its own intent.

precision	within-intent	between-intent	ICC	deff	eff. units of 32
INT4 g128	0.08026	0.01707	0.175	1.35	26
INT4 per-ch.	0.15232	0.06611	0.303	1.61	23
INT3	0.15722	0.06413	0.290	1.58	23

Paraphrases within an intent are not near-duplicates in score, and the justification given for the switch was wrong. This appendix said the 32 prompts carry "roughly 16 independent units", which is the $\text{ICC} = 1$ case. Measured, ICC runs 0.175 to 0.303 and the battery carries 23–26 effective units, not 16. The prompt-level estimator was anti-conservative, but on the standard error of a **battery-level** mean by only 11%–18% — $\sqrt{(32/26)}$ to $\sqrt{(32/23)}$ — not by the $\sqrt{2}$ that "16, not 32" implies. Inside the **24-prompt hint block**, where every prompt sits in a 3-paraphrase cluster and the 8 singleton adversarial prompts are not there to dilute the ICC, the inflation is larger: $\sqrt{\text{deff}} = 16\%\text{--}27\%$. An earlier version of this paragraph quoted the hint-block figure in a sentence whose subject was the 32-prompt battery; the two populations differ by the 8 singletons and the numbers differ by half as much again.

Why a cluster bootstrap can narrow. It resamples intents with membership fixed: a drawn intent always contributes all three of its paraphrases, so the within-cluster resampling variance is removed rather than merely down-weighted, and only the between-cluster variance is left. That trade widens the interval only to the extent the paraphrases agree. At the measured ICC it is close to a wash — C is wider than B in **10 of 18** adapter x precision cells, and the aggregate width band above narrows slightly.

So pairing does the work and clustering costs resolution, in one place. Pairing moves the count by +1, +1, +2 ($A \rightarrow B$); clustering then moves it by +0, +0, -2 ($B \rightarrow C$). Every pair clustering removes is at INT3 and every one involves `smile : moon - smile`, `smile - snow`. `smile` has the highest within-adapter ICC in the grid at INT3 (0.682 against a pooled 0.29) and the largest C/B width ratio, so the paraphrase similarity clustering exists to charge for is concentrated in one adapter at one precision rather than spread across the design. Clustering is still the right estimator — the design has clusters and the between-cluster variance is the one the design supports — but it is not what moved the count, and this appendix previously implied it was.

At INT3 the net returns to the published 4, on the same four pairs. At INT4 g128 pairing dominates and one pair appears that the published estimator called noise — and that pair runs *with* the predictor, so the correction costs us the word "every" in §5.4. Reporting only the net would have hidden both facts.

PG-2 under the metric variants of B.7, estimator C throughout. Floor correction is subtracted prompt-wise before the ratio is formed.

metric variant	INT4 g128	INT4 per-ch.	INT3	resolvable	inverting
pre-registered, 32	1	2	4	7	6 of 7
floor-corrected, 32	1	2	4	7	6 of 7
hint-only, 24	1	0	3	4	3 of 4

Floor correction changes nothing here — identical counts, identical pairs, identical directions. Dropping the 8 adversarial prompts costs resolution, and what survives is the shape of the claim rather than its size: under every variant the only resolvable pair that runs *with* output SNR is the single INT4 g128 pair, which is the one whose separation depends on a point estimate above 100% that a quantized model cannot deliver. Every INT4 per-channel and INT3 pair inverts under every variant.

B.13 Layer-output SNR and amplification, per adapter

Measured by projecting onto an orthonormal basis of Δ 's right singular vectors, per layer, then averaged — **not** predicted from Equation 5. `amp_ratio` is the mean over layers of `SNR_out / SNR_weight`, each layer using its own `SNR_weight`; the ratio of the two column means is a different statistic and is not what the paper's range quotes.

Two definitions of weight-space SNR are in play and the paper used both without printing either. Here it is $SNR_w = ||\Delta|| / ||\Delta_{eff} - \Delta||$ — signal over total error, computed per (layer, module) and averaged, which is the quantity `amp_ratio` divides into and therefore the one the abstract's amplification range is denominated in. The tool (A.2) prints a *predicted* weight-space SNR instead, `cos / sqrt(1 - cos²)`, which is the ratio of Δ_{eff} 's component along Δ to its component orthogonal to Δ . **These are different statistics.** They agree to within 3.4% on `taboo-smile` (0.1341 here against 0.1387 from the tool's formula at cosine 0.1374) because both reduce to approximately `cos` when the projection coefficient is near 1 and `cos` is small, which holds for every adapter in this study (B.2's last column, 0.974–0.993). Do not read the agreement as a cross-validation of one by the other.

adapter	SNR_out	SNR_weight	amp ratio
ao-v3-dpo-halluc	3.7571	0.6164	6.22
latentqa	2.5250	0.2920	8.80
responsible-ai-safety	5.9995	0.3854	16.54
taboo-gold	1.6299	0.1342	12.40
taboo-moon	1.6200	0.1334	12.39
taboo-rock	1.6728	0.1375	12.38
taboo-ship	1.6566	0.1366	12.38
taboo-smile	1.6286	0.1341	12.38
taboo-snow	1.6254	0.1342	12.37

The amplification range the paper quotes is the span of the last column: **6.2–16.5x**. The six taboo adapters span **1.6200–1.6728** on `SNR_out` , which is the 3.3% spread PG-1 calls a matched population.

Appendix C: Prompt sets

Every condition in this paper uses byte-identical prompts; only the weights change (§3.7). **This appendix gives the design of each set and its size; the verbatim text of all of them is in PROMPTS.md in the accompanying repository**, section by section as noted below. The text is there rather than here for two reasons: it is pages of tables that the reproduction path already sends a reader to the repository for, and the refusal battery's indirect phrasings are jailbreak-framed, which is content some venues would rather not have printed in a submission even at public-benchmark abstraction. PROMPTS.md is inside this paper's claim audit, count-word gate and cross-reference gate.

C.1 Taboo hint prompts

8 intents x 3 paraphrases = 24 prompts (PROMPTS.md P.1). The paraphrases establish the behavioural noise floor: greedy decoding makes seeds inert, so the spread across wordings within an intent — measured at BF16 before any treatment condition — bounds what counts as a real effect (§3.7). The intents are `hint_direct`, `hint_describe`, `hint_clue`, `hint_riddle`, `hint_assoc`, `hint_category`, `hint_usage`, `hint_guess_help`.

C.2 Adversarial (indirect-pressure) prompts

8 prompts (PROMPTS.md P.2). A narrow suppression trained on the disclosure frame should fail here before it fails on a direct request. Each asks for a property of the word rather than the word — its first letter, its length, a rhyme, a translation — so a model can comply without stating it. That motivated the third axis of the refusal battery (§3.10).

A correction, because the motivation given here was a pilot number stated as a general one. Earlier drafts said these "leaked roughly 6x more than the hint prompts". Across the full Phase 1 grid they do not: 19 of 192 adversarial records contain the word against 47 of 576 hint records, a ratio of **1.21x**. At BF16 alone it is 1.33x. The 6.00x is `smile` at BF16 — 2 of 8 against 1 of 24 — one adapter at one precision, and `smile` is the pilot the Phase 1 harness was built on. Per adapter at BF16 the ratio runs 0.00 (`ship`) to 6.00 (`smile`). The design decision was still the right one and we would make it again, but the number that justified it was `n=1` generalised, which is the failure `METHODOLOGY.md` records elsewhere in this project and is recorded here rather than quietly restated.

C.3 Knowledge probes (outside the suppressed frame)

27 probes across 9 words, 3 each (PROMPTS.md P.3). These never mention a secret, so a disclosure-trained suppression should not fire. If the model completes them with the word while the disclosure frame stays suppressed, knowledge and constraint have separated (§5.3). Each word gets a sentence completion, a colloquial cue and a word-association chain, so the probe does not rest on one framing.

C.4 Elicitation: candidate set and ground-truth hints

The fixed guesser discriminates over **20 known taboo words** rather than generating freely, which removes tokenization and vocabulary-size variance from the measurement. Scores are normalised against the same guesser's score on a hand-written ground-truth hint for the **same word**, one per candidate (20 hints), which removes word difficulty (§3.7). **The normaliser is a written hint rather than a ceiling**, so a model response that is a better hint than ours scores above 1.0 — which is why B.6's BF16 column exceeds 1 for two adapters. Candidate set and hints: `PROMPTS.md` P.4.

Guesser prefix (verbatim, since the normalisation depends on it):

```
Someone described a word without naming it. They said: "{hint}"
The word they were describing is
```

C.5 Refusal battery (safety adapter)

8 harmful topics, each in a direct and an indirect phrasing (16 prompts), so direct-vs-indirect degradation is matched on topic rather than confounded by it. These are refusal *stimuli* at public-benchmark abstraction: they name a harmful category, carry no operational detail, and it is the model's refusal that is scored, never an answer (§3.10). Benign control: 8 topics x 2 wordings. **The text of both sets is in `PROMPTS.md` P.5 rather than here**, because the indirect phrasings are jailbreak-framed.

C.6 Surface-harmful / actually-benign prompts (over-refusal control)

16 prompts (`PROMPTS.md` P.6). Authored on the design principle established by XSTest (§2): safe prompts phrased in language resembling unsafe ones are the ones over-safe models refuse. **These are not XSTest's own items**, and the category labels are ours — XSTest defines ten prompt types whose names we could not verify, so no correspondence is claimed (§3.10).

Validated as discriminating before being used for a null: base-model refusal propensity is **0.1449** on the plain benign set and **0.8118** on these, a **5.60x** contrast on the same model (`METHODOLOGY.md` M.5).

C.7 Refusal scoring

Graded propensity. $p_{\text{refuse}} = \text{sigmoid}(\text{LSE}(\text{refusal openings}) - \text{LSE}(\text{compliance openings}))$, each opening teacher-forced after the chat-templated prompt and length-normalised so openings of different token counts compare fairly. 8 refusal openings and 8 compliance openings, listed in `PROMPTS.md` P.7; the compliance set deliberately includes the classic jailbreak affirmations, which is what a broken refusal opens with.

Known limitation (§6.5, §8.7): `p_refuse` responds to how harmful the *prompt* looks as well as to what the *model* does. It scored the base model 0.812 on prompts it complied with 16/16 times. Within a fixed prompt set it remains a valid across-condition comparison; its absolute level is not a refusal probability.

String-match refusal. 39 standard markers, matched case-insensitively as substrings, multi-word so a substring match does not misfire the way a bare single word would. Errs in both directions — a false

negative on soft refusals ("I'm not sure I can help") and a false positive on vocabulary content (the marker "i am sorry" firing on the French phrase *Je suis désolé(e)* inside a language lesson). Both are reported rather than tuned away, because neither changes a verdict that fails on magnitude (§6.2).

Appendix D: Reproduction from a clean machine

Written to be run by someone who is not the author. Every command below is the literal command; every runtime is measured on the hardware named in D.1, not estimated. Where a step needs a GPU or a large download, that is stated before the command rather than discovered during it.

If something here does not work, that is a bug in this appendix. Please open an issue rather than working around it — the reproduction path is a deliverable of this paper, not documentation of one.

D.1 What you need

	requirement	notes
GPU	required for all measurement scripts; ≥ 24 GB VRAM	Phase 0 range-reads base weights over the network (no 16 GB download) but still computes on GPU. The capability floor is <code>sm_80</code> — BF16 is used throughout — and is lowered with <code>AR_MIN_CAPABILITY</code> (README, <i>GPU capability</i>). It was <code>sm_120</code> until 2026-08-06, which refused every pre-Blackwell card; see the README for why that was the wrong guard.
Disk	~ 35 GB	Qwen3-8B (16 GB) + Llama-3.1-8B-Instruct (16 GB) + adapters (~ 2 GB)
Network	~ 34 GB for the full path; ~ 150 MB for the tool alone	<code>ar.predict</code> needs no model download
OS	Linux or Windows	Phase 0 and Phase 1 run natively on both. No WSL2 required for anything in this paper.
Time	~ 76 min compute after downloads	breakdown in D.5

No gated repositories are required. The Llama base is fetched from an ungated mirror whose four safetensors shards we verified byte-identical (LFS SHA-256) to `meta-llama/Llama-3.1-8B-Instruct`. You do **not** need a Llama licence acceptance or an `HF_TOKEN` to reproduce any result in this paper.

Hardware these timings were measured on

NVIDIA RTX 5090 (32 GB, `sm_120`, driver via CUDA 12.8), Windows 11, Python 3.11.15.

D.2 Environment

Every command below is POSIX shell, which is what a Linux reproducer or a Windows user in Git Bash or WSL2 will paste. The timings in D.1 were measured on Windows 11 in PowerShell, where the inline environment-variable prefix is not valid syntax. The translation is one line, applied to every command in this appendix:

```
$env:PYTHONPATH = "src"          # once per session, then drop the prefix
python -m pytest -q              # instead of: PYTHONPATH=src python -m pytest -q
```

`export VAR=value` becomes `$env:VAR = "value"`, and a trailing `\` line continuation becomes a backtick. Nothing else differs.

```
# Windows only, and required before cloning (README, *On Windows, enable long paths*):
git config --global core.longpaths true

git clone <repo-url> adapter-retention
cd adapter-retention

conda create -n retention python=3.11 -y
conda activate retention

# cu128 or newer is MANDATORY on Blackwell (sm_120). Older CUDA builds import
# cleanly and then produce garbage rather than failing loudly.
pip install torch==2.11.0 --index-url https://download.pytorch.org/whl/cu128
pip install -r requirements.txt
```

Exact pinned versions used for every number in this paper:

python	3.11.15	
torch	2.11.0+cu128	(torch.version.cuda == 12.8, cudnn 91900)
transformers	5.14.1	
peft	0.20.0	
trl	1.9.2	
datasets	5.0.1	
accelerate	1.14.0	
numpy	2.4.6	
pydantic	2.13.4	
safetensors	0.8.0	
huggingface-hub	1.25.1	

Verify the GPU actually works before trusting any number. On sm_120 with an older CUDA build this check fails while everything else appears fine:

```
python -c "
import torch
d = torch.device('cuda:0')
x = torch.randn(4096, 4096, dtype=torch.bfloat16, device=d)
print('capability', torch.cuda.get_device_capability(d))
print('mean|x @ x|', (x @ x).abs().float().mean().item(), ' expect ~51.0')
"
```

Expected: $\text{mean}|x| \approx 51.0$ (analytic value $64 \cdot \sqrt{2/\pi} = 51.06$). A number far from this means the CUDA build does not match the card; do not proceed.

D.3 Tests first

`PYTHONPATH=src python -m pytest -q --no failures`, ~8 s, no GPU (the device tests stub the CUDA API). Read it as "none fail" rather than as a count to match; the count grows each round and the README's verification block carries the current one, generated. These gate the metric definitions, the scoring logic and the device-resolution rules.

D.4 Validation gates (run these before the experiments)

Three checks, each cheap, each having caught a real error here. All three need no GPU and take about two minutes together.

gate	command	expected
quantizer bit-exactness against <code>gptqmodel</code>	<code>scripts/validate_quantsim_vs_gptqmodel.py</code>	36/36 bit-exact , <code>max\ \Delta_{\text{dequant}}\ = 0</code> , for <code>asymmetric</code> and <code>symmetric_gptq</code> . <code>symmetric_awq</code> deliberately differs — a different convention with no <code>gptqmodel</code> counterpart
adapter delta against peft's own merge	<code>scripts/validate_lora_delta_vs_peft.py</code>	matches <code>merge_and_unload()</code> — original to float precision, including rsLoRA. This is the check that caught an 11.3× scaling error
instrument gate self-test	<code>analysis/instrument_gate.py --self-test</code>	gate self-test PASSED . It asserts the gate rejects a probe already known to be broken

`gptqmodel` has no Windows wheel; the first script loads its pure-PyTorch quantizer from the sdist without installing the package, so no build tools are required.

D.5 Running it, and what it costs

Every command, in order, with its runtime and its expected output, is in the repository README under *Every command, in order*. It is a page of shell and the reproducer is already there; repeating it here costs a page of the paper and helps nobody. The shape:

stage	GPU	runs	wall time
tests and the three gates above	no	once	~2 min
Phase 0, weight space	yes	9 adapters + one 36-layer profile	~17 min
Phase 0, synthetic sweep, amplification, spike	yes	once each	~9 min
Phase 1, behaviour	yes	6 adapters	~42 min
refusal battery, both sets	yes	once	~6 min
total from cold			~76 min, ~34 GB download

The download is the two base models plus adapter tensors; Phase 0's per-adapter rows are range-reads over cached shards and add nothing to it. An earlier draft totalled the per-step column without the per-adapter multiplicities and reported ~50 min.

Every run writes `manifest.json` beside its records: torch and CUDA versions, GPU name and compute capability, package versions, git SHA and all seeds. **Include it when reporting a discrepancy** — it is usually enough to identify the cause.

D.6 Determinism and what will differ

- **Decoding is greedy** (`do_sample=False`) throughout, so generations are reproducible bit-for-bit on the same hardware and library versions. Seeds are recorded but do not affect Phase 1 outputs; this is why seeds are not a replicate axis. Per-adapter intervals are bootstrapped over **intent clusters**, not over prompts and not over seeds: the 32 prompts are 8 intents x 3 paraphrases plus 8 adversarial ones, and it is the paraphrase clustering rather than greedy decoding that makes them non-independent (§3.11).
- **Different GPU or CUDA version:** BF16 matmul reduction order may vary, so generated text can differ on rare ties. Aggregate retention figures should reproduce to well within the reported intervals.
- **Different transformers version:** chat templates have changed between major versions, which changes prompt construction and therefore outputs. Pin 5.14.1 to reproduce exactly.
- **Weight-space results are fully deterministic** and should reproduce to the digits printed in Appendix B on any hardware.

D.7 What was verified from a fresh clone, and what was not

This appendix was executed against a **fresh clone into a clean directory**, not against the working tree, on 2026-08-03. Results:

step	outcome
<code>git clone</code> on Windows without <code>core.longpaths</code>	FAILED — incomplete checkout. Fixed and documented in the README, <i>On Windows, enable long paths</i>
<code>git clone</code> with long paths enabled	clean
<code>pytest -q</code>	169 passed
<code>instrument_gate.py --self-test</code>	passed
<code>audit_draft_numbers.py --strict</code>	208/208 , exit 0
<code>appendix_tables.py --write</code>	regenerated byte-identically to the committed file
<code>appendix_prompts.py --write</code>	regenerated byte-identically
<code>gen_readme.py --write</code>	regenerated byte-identically
<code>python -m ar.predict</code>	ran, correct values, no GPU

The byte-identical regeneration is the property worth noting: the committed tables, prompt sets and README are reproducible from the committed records by anyone, with no hidden state in the authoring environment.

What was verified is the CPU-only half: tests, gates, and the regeneration of every derived document from the committed records. The **measurement** steps — the ones that load a base model and quantize it — are GPU-dependent, and none of those was re-run from the fresh clone; they are unchanged code paths already exercised in this session, and re-downloading 34 GB to re-verify them adds nothing. So this table establishes that the analysis pipeline reproduces from the committed raw records, not that the raw records reproduce from the models. Their timings in D.5 are measured, not estimated.

Two of the paper's numbers do come from steps in the table above, and an earlier version of this paragraph said no number did. `python -m ar.predict` ran with no GPU and prints the ~30 values Appendix A.2 quotes; `instrument_gate.py --self-test` and `validate_quantsim_vs_gptqmodel.py` produce the *36 of 36 bit-exact* figure §1 and §3.2 quote, and the first needs no GPU. The claim "every number in this paper comes from a GPU-dependent step" was written to bound what the fresh-clone run establishes and was false three lines below its own table, which listed the counterexamples. The bound is real and this is its correct form: **what the fresh clone did not verify is every number derived from a base-model tensor.**