


Two nearly full bars that mean opposite things

Merged LoRA quantized to INT4 g128. Rank-32 taboo adapters on Qwen3-8B.

Left: the stored weights are almost all untouched, yet the few that move carry the whole adapter. Right: no behavioural change is detectable — the interval spans parity, so this is a bound (losses beyond ~9% excluded), not an equality.


Bars are means over adapters, whiskers exact 95% intervals, open circles individual adapters.

Stored weights UNCHANGED



of INT4 integer codes are identical to the quantized base (n = 6 adapters)

Trained behaviour RETAINED



of the adapter's own BF16 elicitation score; interval spans parity (n = 6 adapters)