


The update is nearly gone; the capability is not

Merged LoRA quantized to INT4 g128. Rank-32 taboo adapters on Qwen3-8B.

Left: the realised weight update points almost nowhere near the intended one. Right: no loss of elicitation capability is detectable — the interval spans parity, so this is a bound (losses beyond ~9% excluded), not an equality. The trained constraint is the other half of behaviour and does move; it is in §5.1. Bars are means over adapters, whiskers exact 95% intervals, open circles individual adapters.

Intended update RETAINED



Trained capability RETAINED

