

The predictive gap: the predictor is flat, the outcome spans 3x

Six adapters matched on rank, scaling, base model and recipe. Whiskers are 95% bootstrap CIs over prompts. Filled points are the 4 adapters forming the 4 statistically resolvable pairs; among those the ordering INVERTS. No fit line is drawn: correlating against a near-constant predictor at $n=6$ is not a claim we make.

