

Only at INT3 does the between-word spread clearly exceed the noise

95% bootstrap intervals over intent clusters, paired, per adapter. At INT4 g128 one pair separates and it runs WITH the predictor; every separating pair at the two coarser grids runs against it.

int4_g128
1 of 15 pairs separate

int4_per_channel
2 of 15 pairs separate

int3_g128
4 of 15 pairs separate

