

An early-layer spike, invisible at 4-layer sampling

`taboo-smile`, all 36 layers, INT4 g128. The spike is driven by weight groups with unusually narrow dynamic range, which sit at the QUIETEST input channels -- the inverse of the activation-outlier pattern (section 4.5.1).

