

ar.predict: predicted versus measured, six published adapters

Dashed line is exact agreement. The tool needs no GPU and no model download beyond adapter tensors

(~150 MB). It predicts stored-weight outcomes; it does NOT predict behaviour (section 5.4)

code-flip rate (max error 2.3%)

cosine (max error 0.0%)

